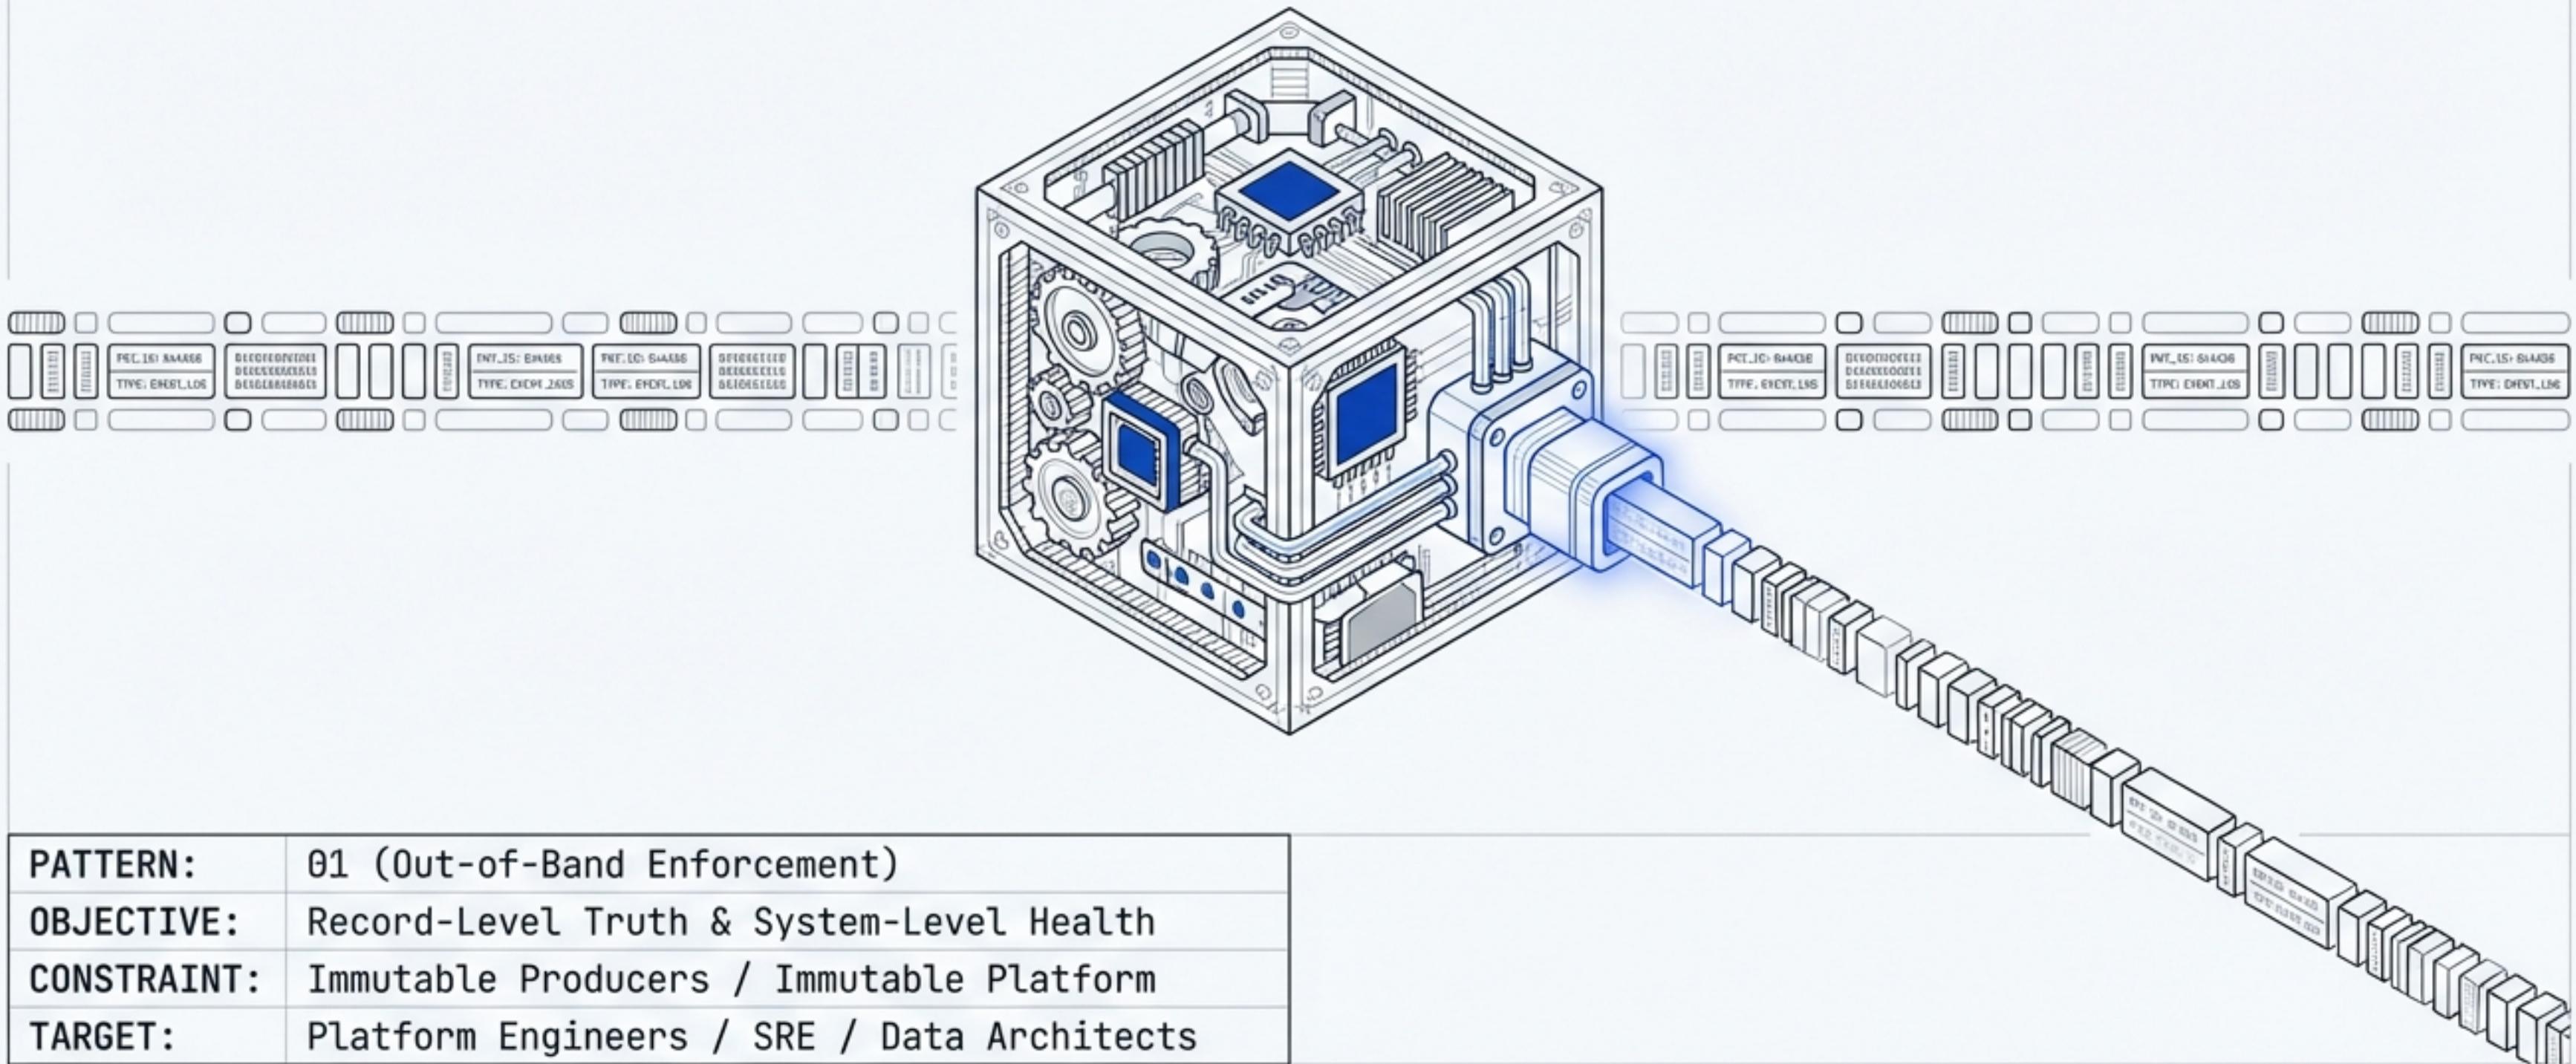


Low-Level Design: Out-of-Band Observability & Policy Enforcement

Establishing Immutable Truth in an Immutable Streaming Environment



The Constraint: Pattern 1 (Out-of-Band Enforcement)

The Challenge:

The Central Streaming Platform and Producers are immutable. We cannot modify their code, change routing logic, or block writes inline.

The Architectural Shift:

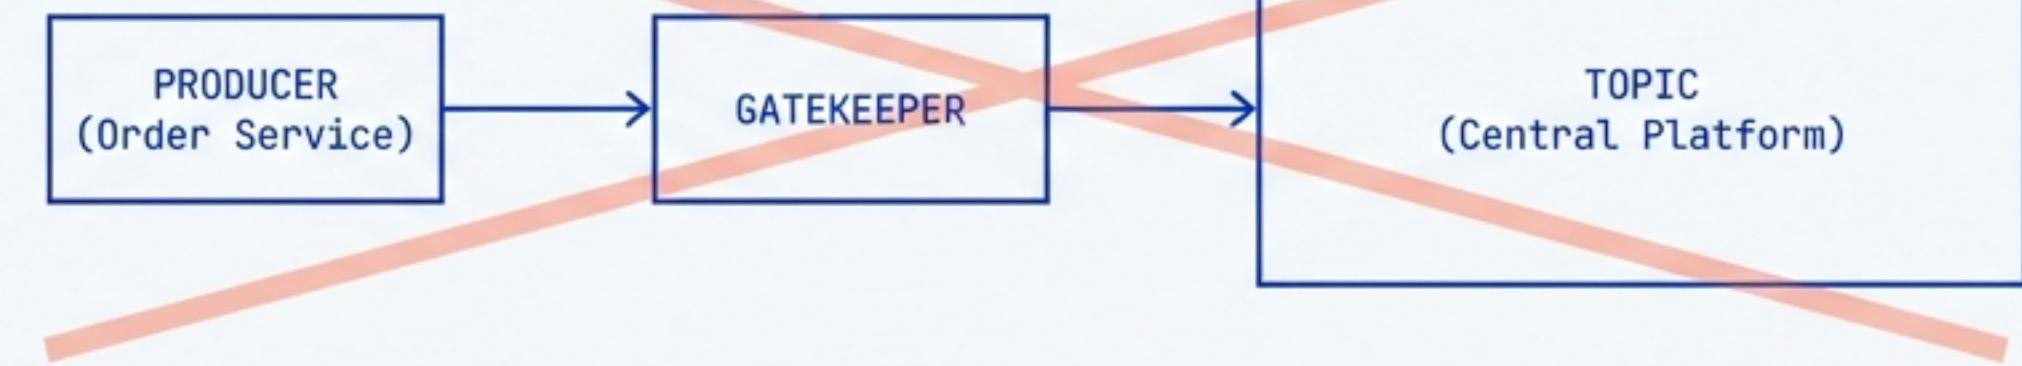
Move from "Gatekeeper" (blocking the write) to "Sentinel" (validating the read).

The Mechanism:

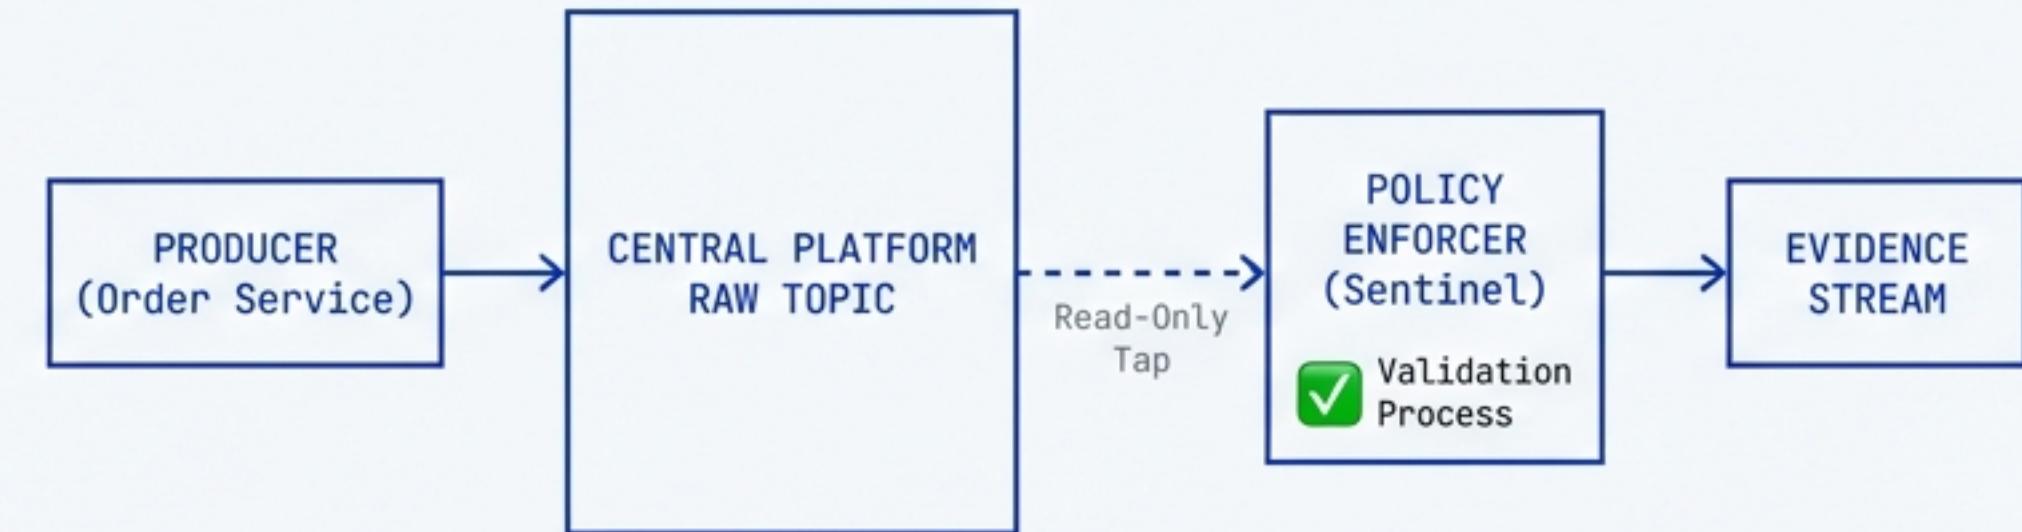
Pre-consumption enforcement. We identify and isolate bad data before it is consumed by downstream applications.

SCHEMATIC DIAGRAM: INLINE VS. OUT-OF-BAND

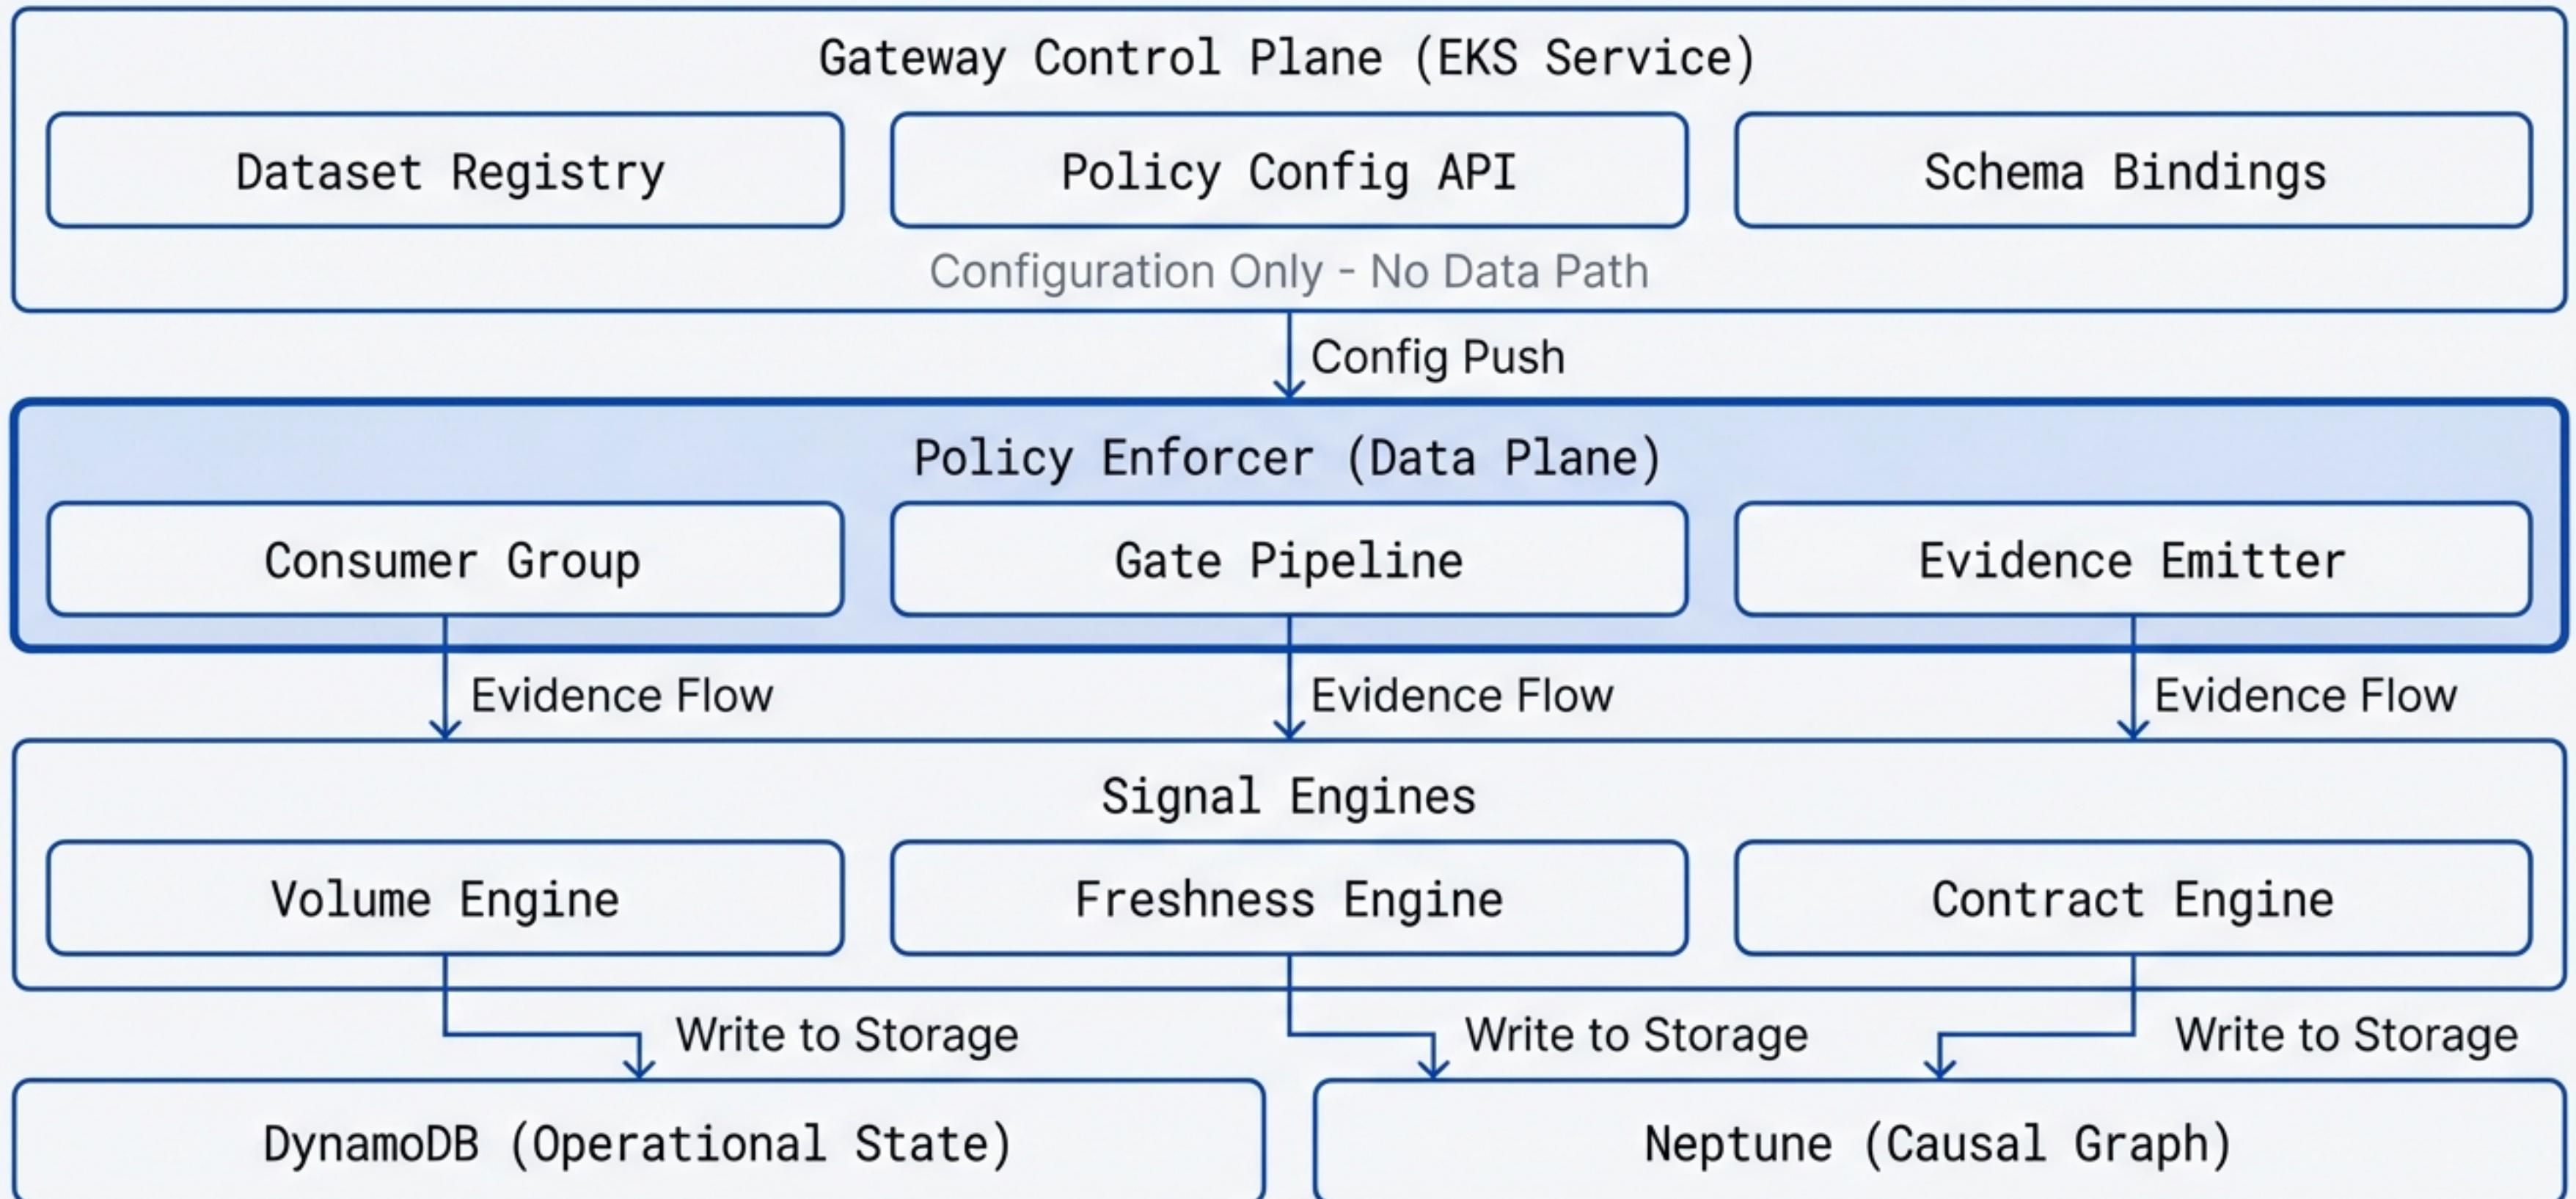
Standard Inline (Not Possible)



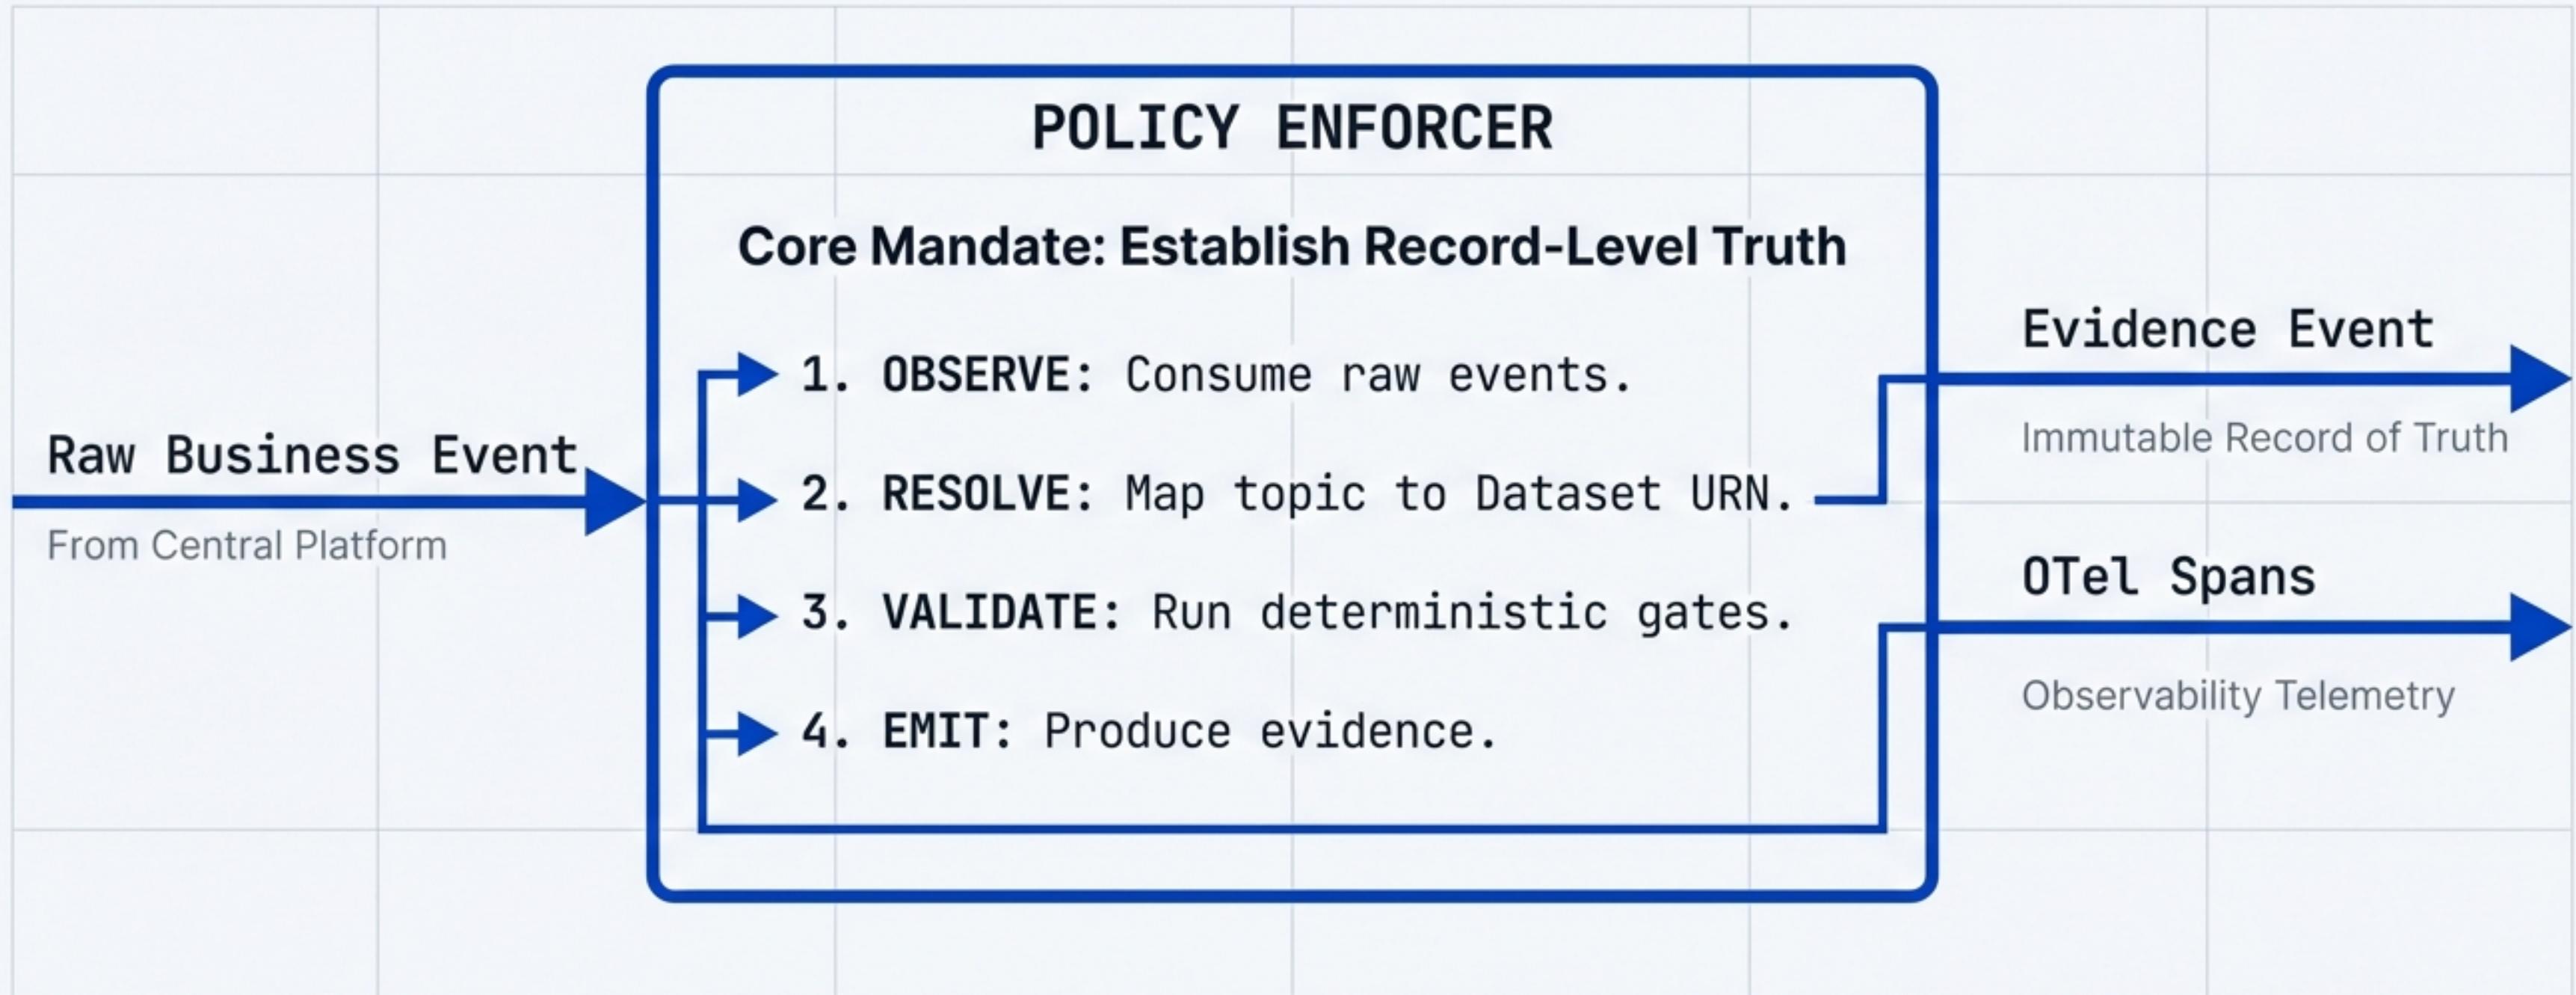
Pattern 1: Out-of-Band (Implemented)



High-Level Component Architecture



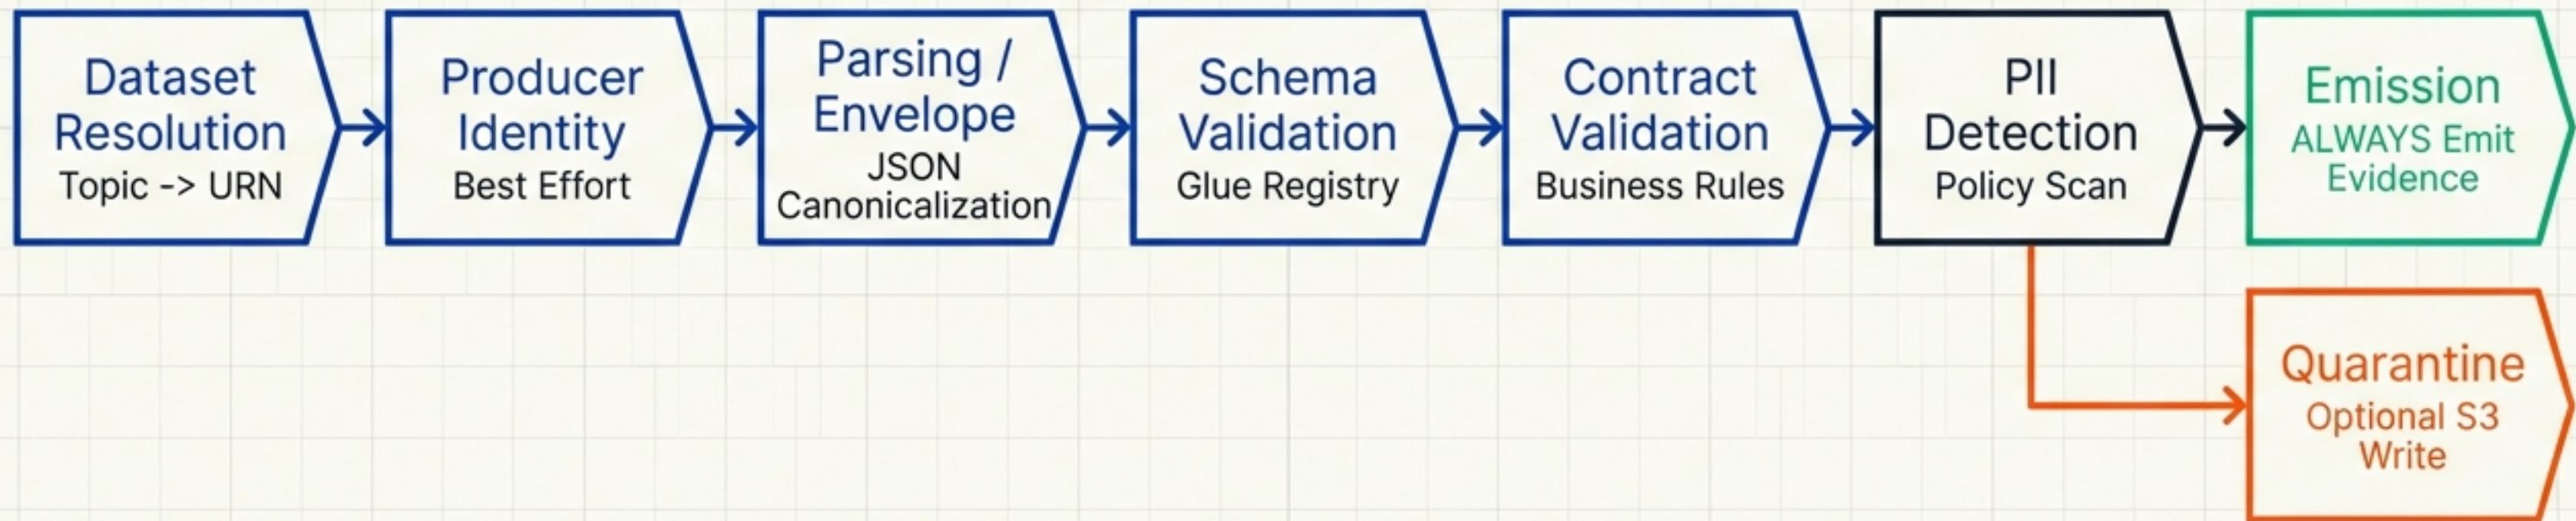
Deep Dive: The Policy Enforcer Responsibilities



HARD BOUNDARIES (What the Enforcer does NOT do):

✗ - No Aggregation (Signal Engine job) ✗ - No Alerting (Incident Mgmt job) ✗ - No Subjective Judgment (Only Facts)

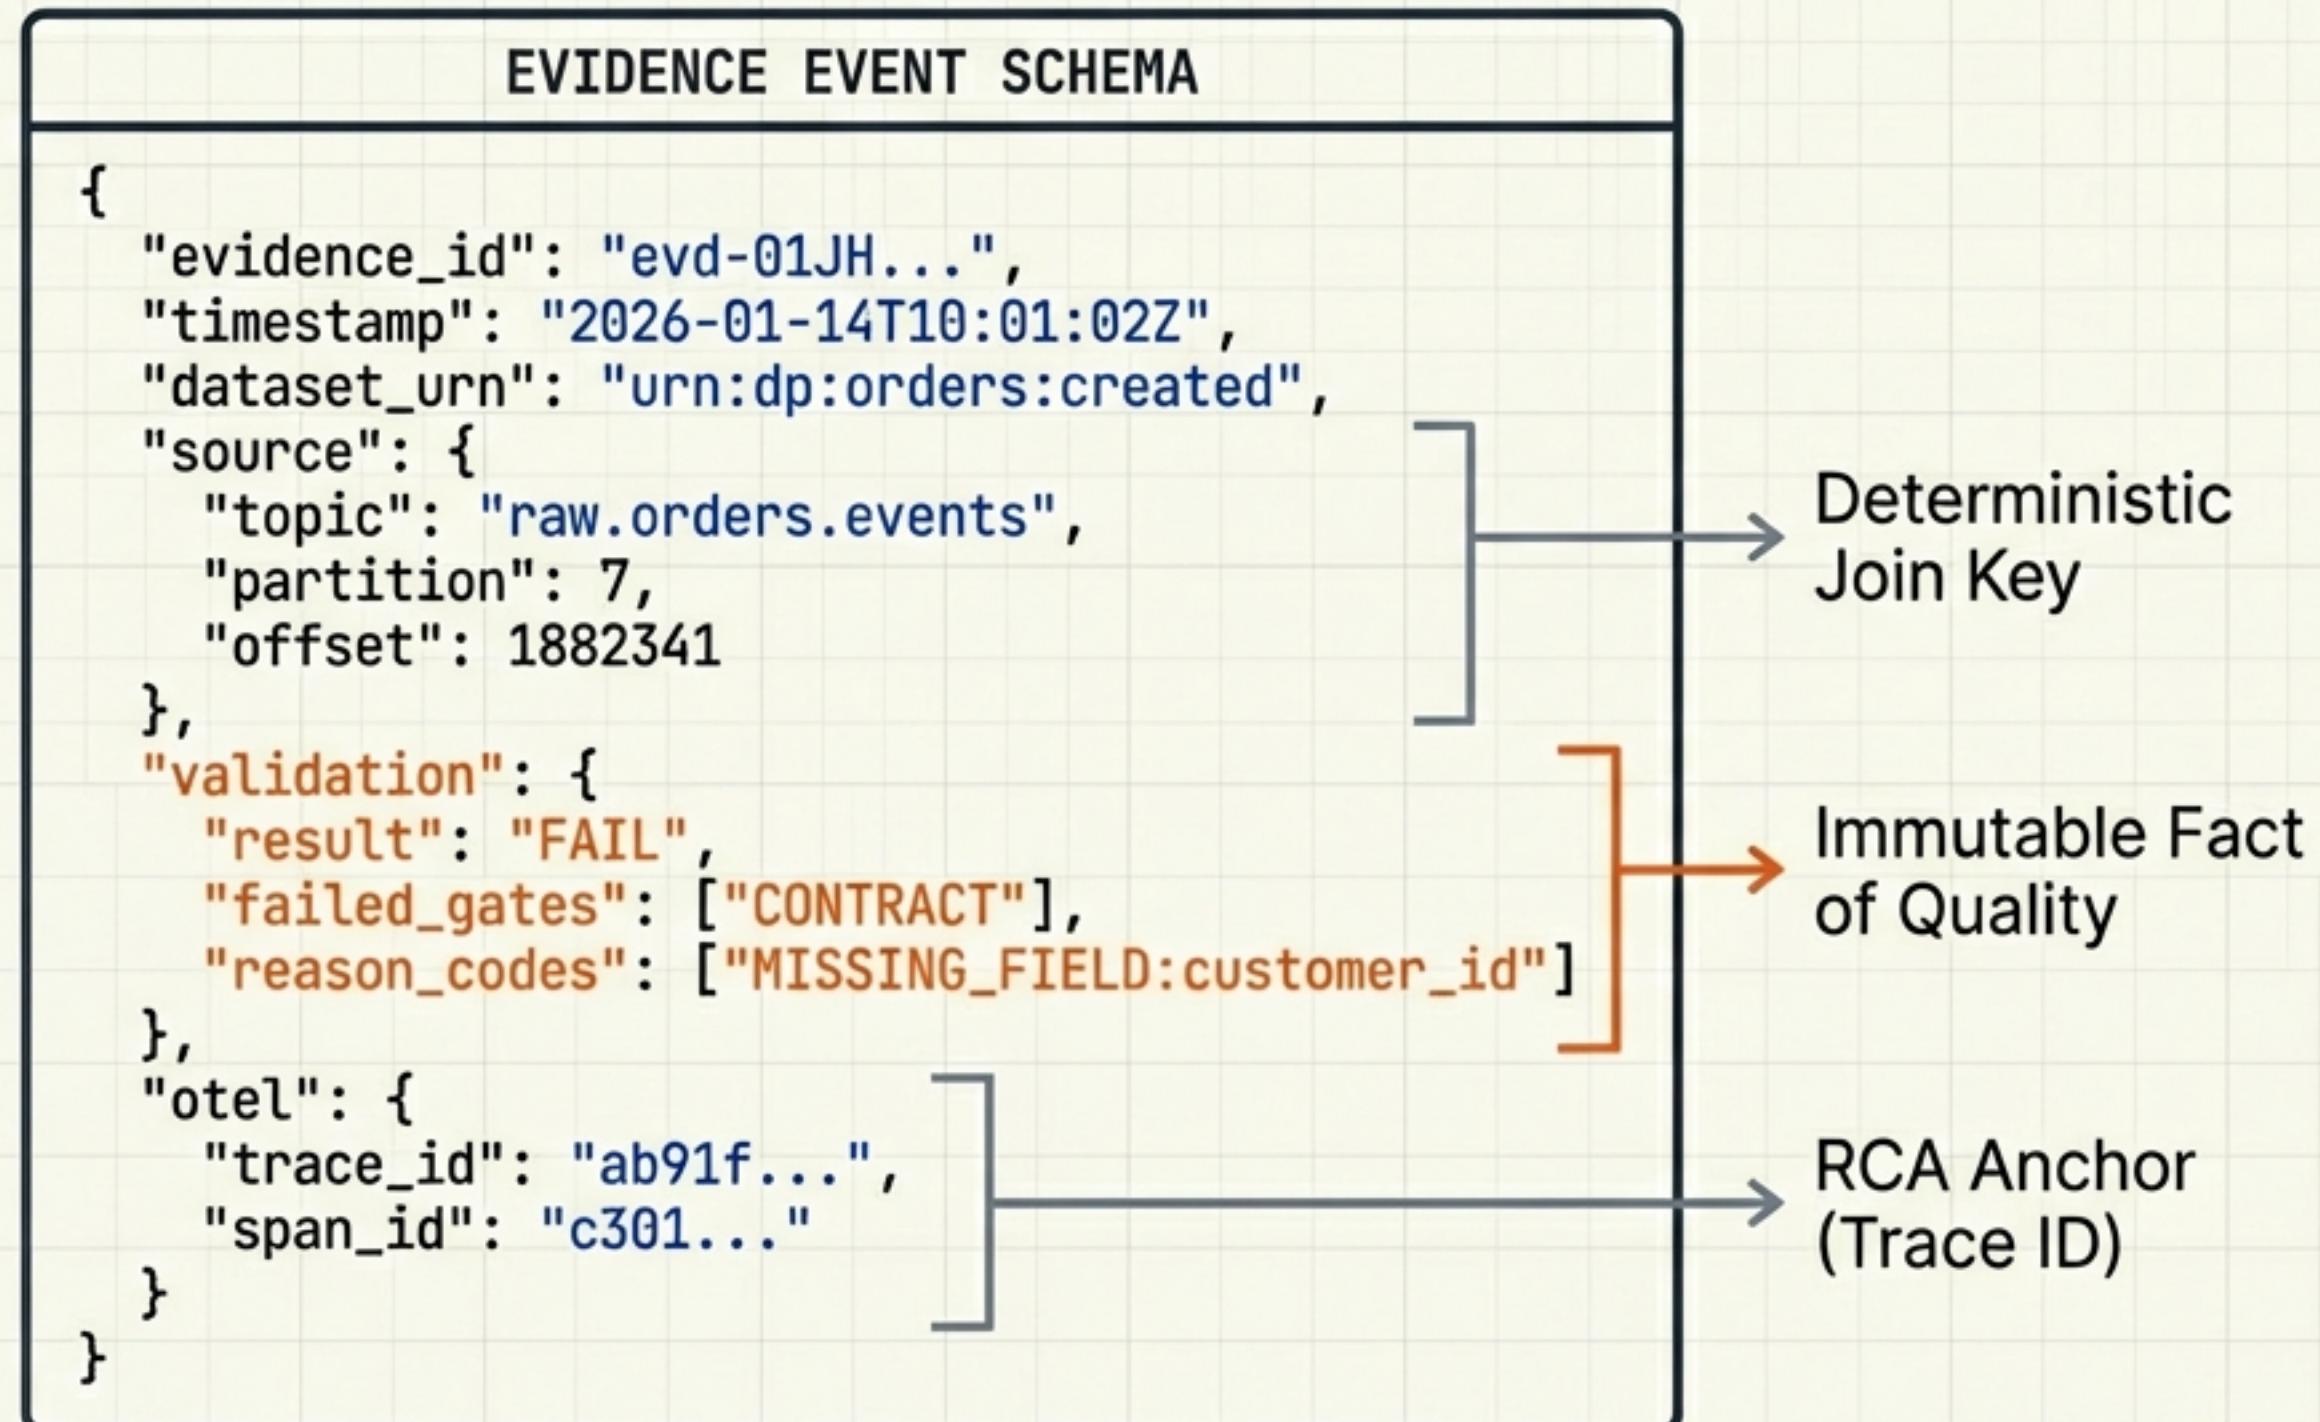
The Gate Pipeline: Deterministic Evaluation



Sequential, deterministic processing for every single record.

The API Contract: Evidence Event Schema

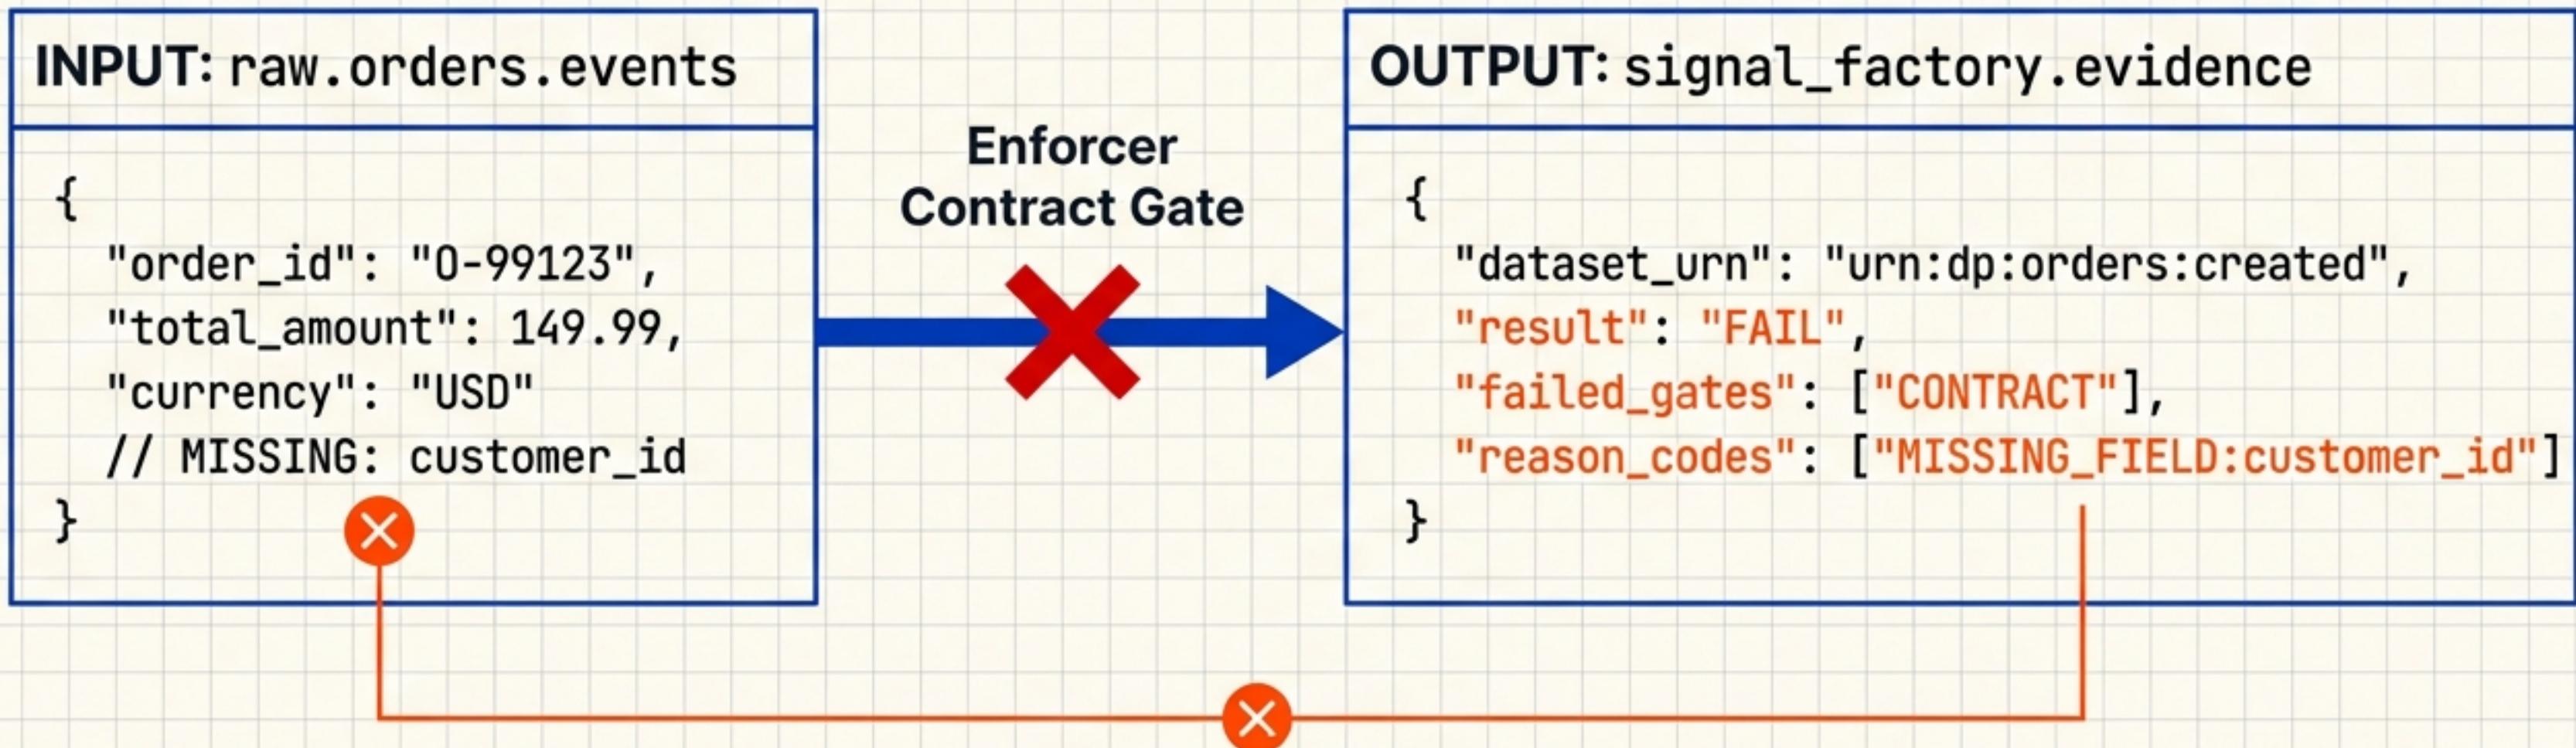
The only stable API between the Enforcer and the rest of the world.



Key Takeaway:
One Evidence
event is emitted for
every raw record,
regardless of
validity. This
creates a complete
audit trail.

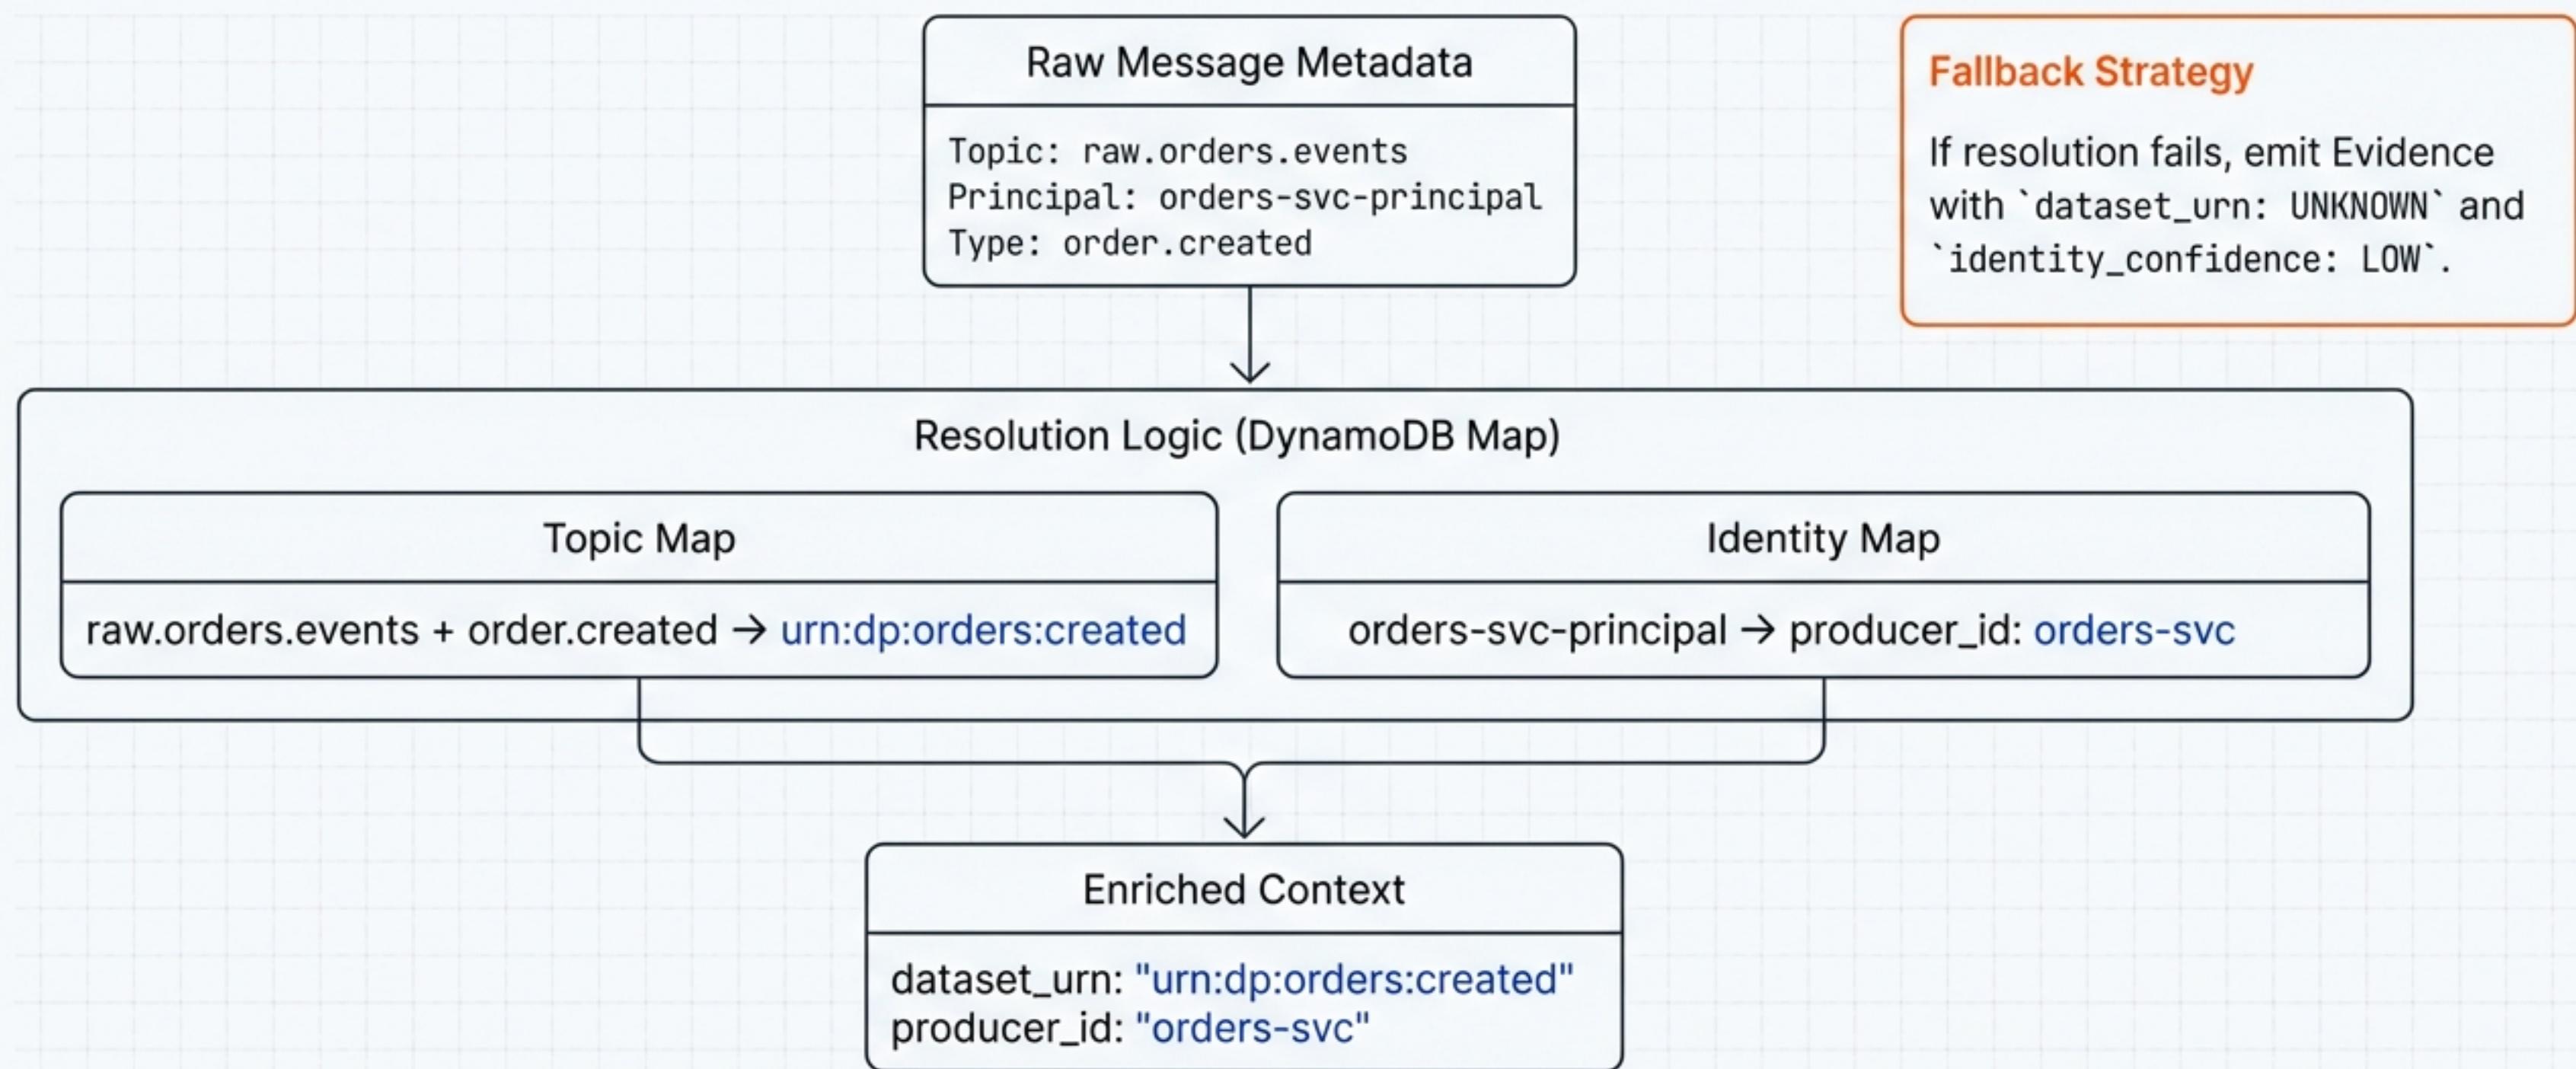
Steel Thread: The 'Order Created' Failure

Scenario: Order Service deploys bad code at 09:58 AM.
The 'customer_id' field is missing.



Dataset Resolution & Identity Inference

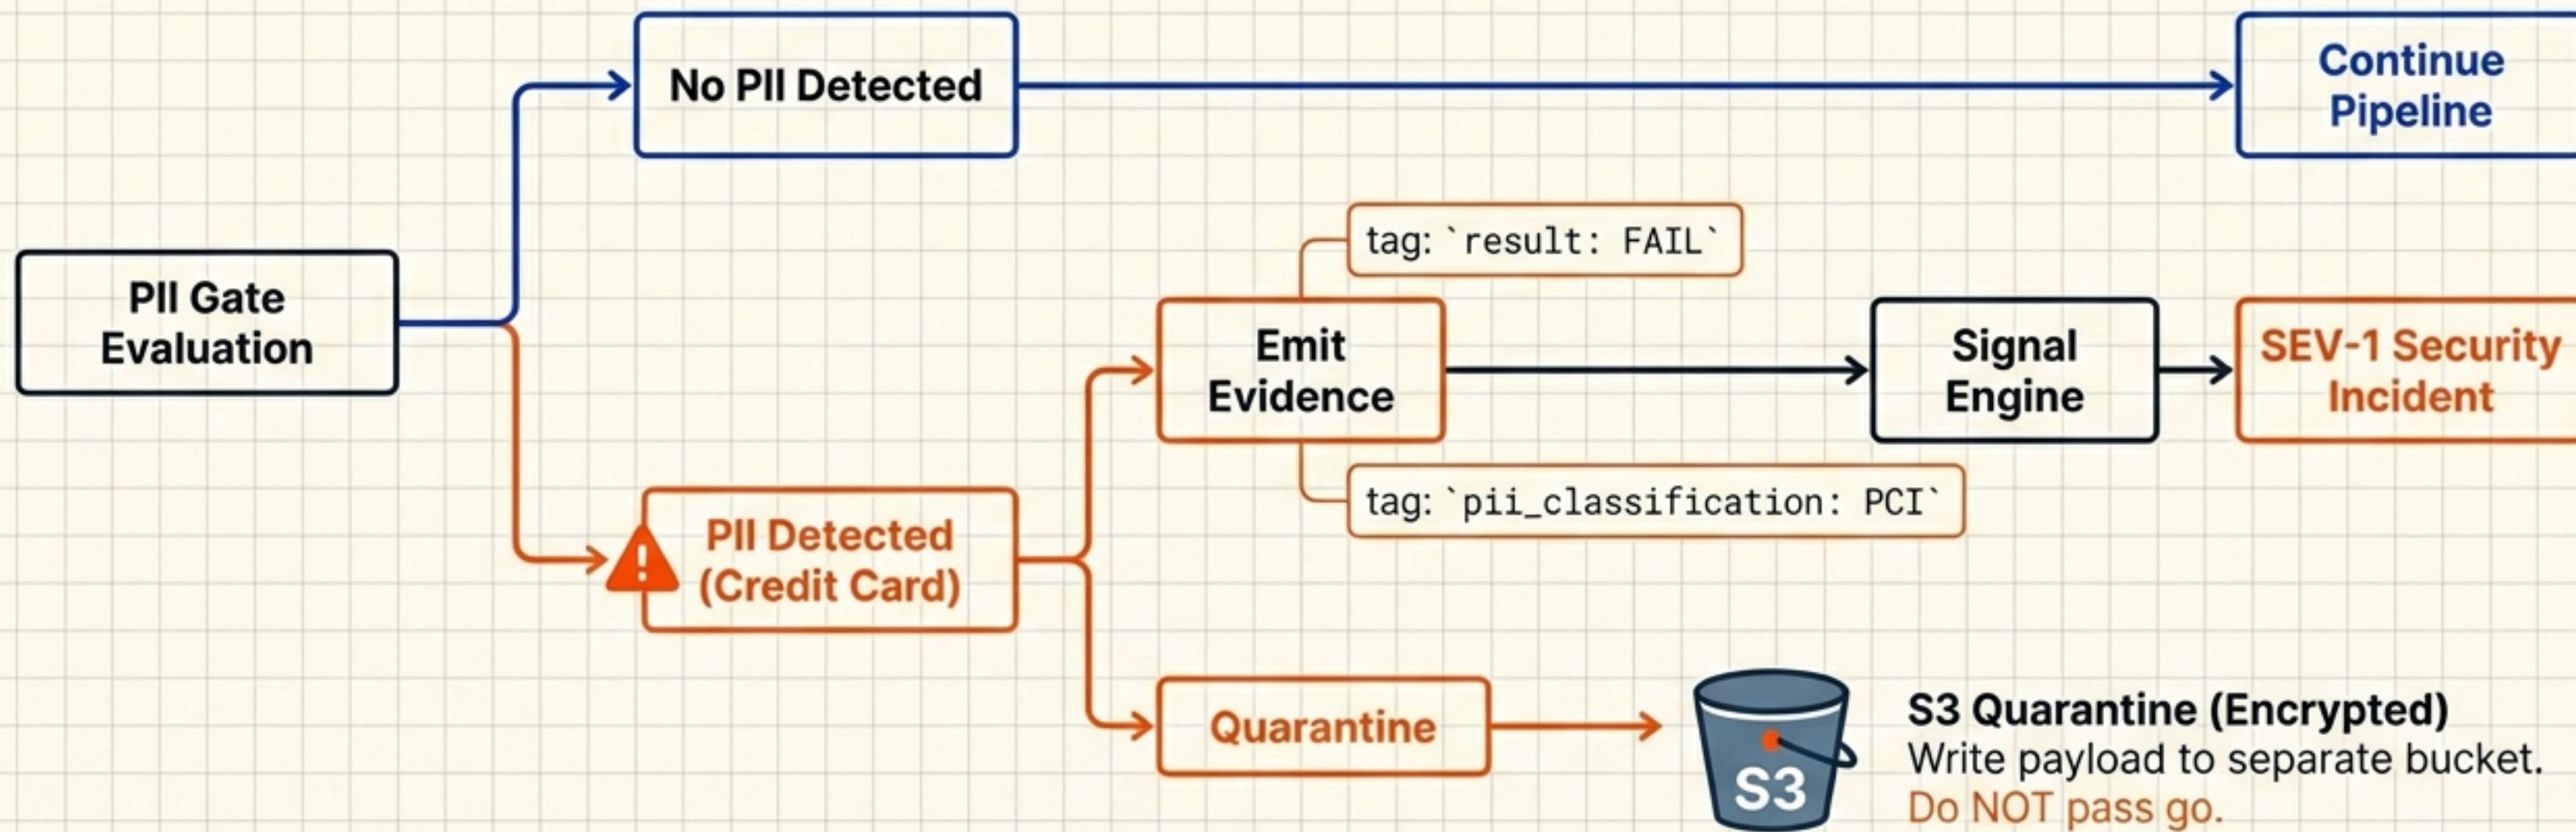
The Challenge: Out-of-band messages don't always carry headers declaring "I belong to Dataset X".



Policy Actions: PII Detection & Quarantine

Scenario: Developer logs a credit card number.

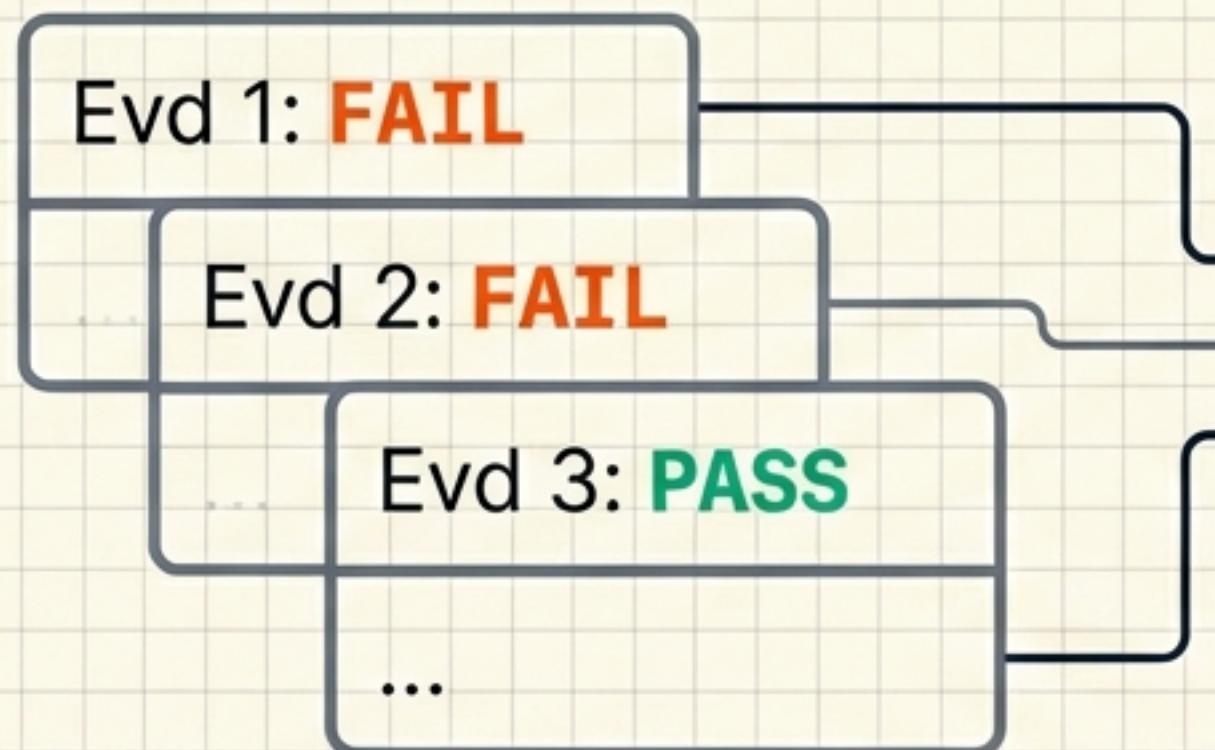
Policy is strict PCI blocking.



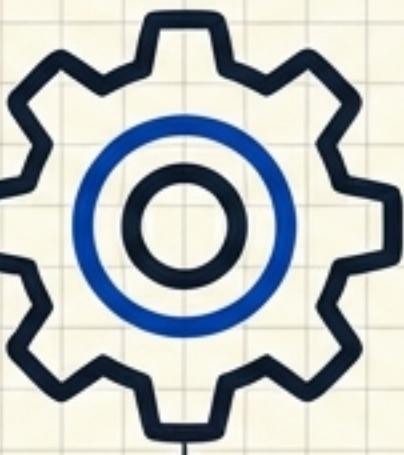
From Record Truth to System Health

Separation of Concerns: The Enforcer creates facts; the Signal Engine creates judgments.

Stream of Evidence Events



Signal Engine (Aggregation)



Window: 5 minutes.
Threshold: 99.5%



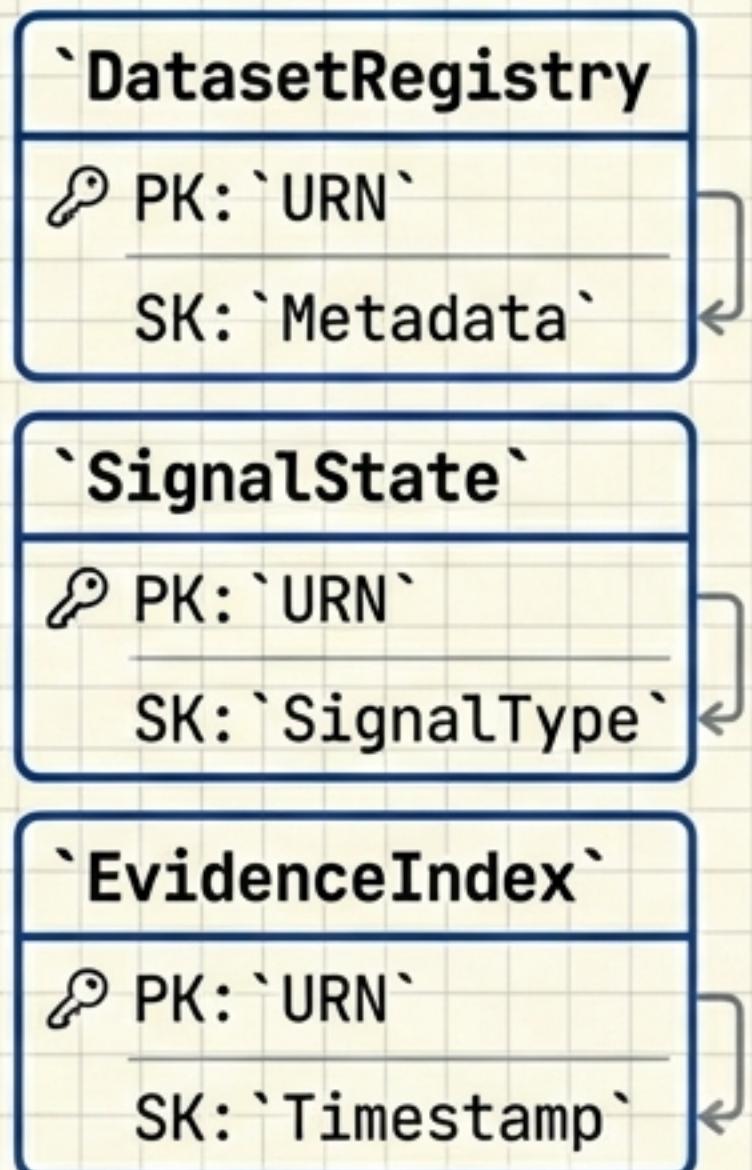
Signal Types

1. **Volume** (Rate drops/spikes)
2. **Freshness** (Time since last valid event)
3. **Contract** (% Compliance)

Data Modeling: State vs. Causality

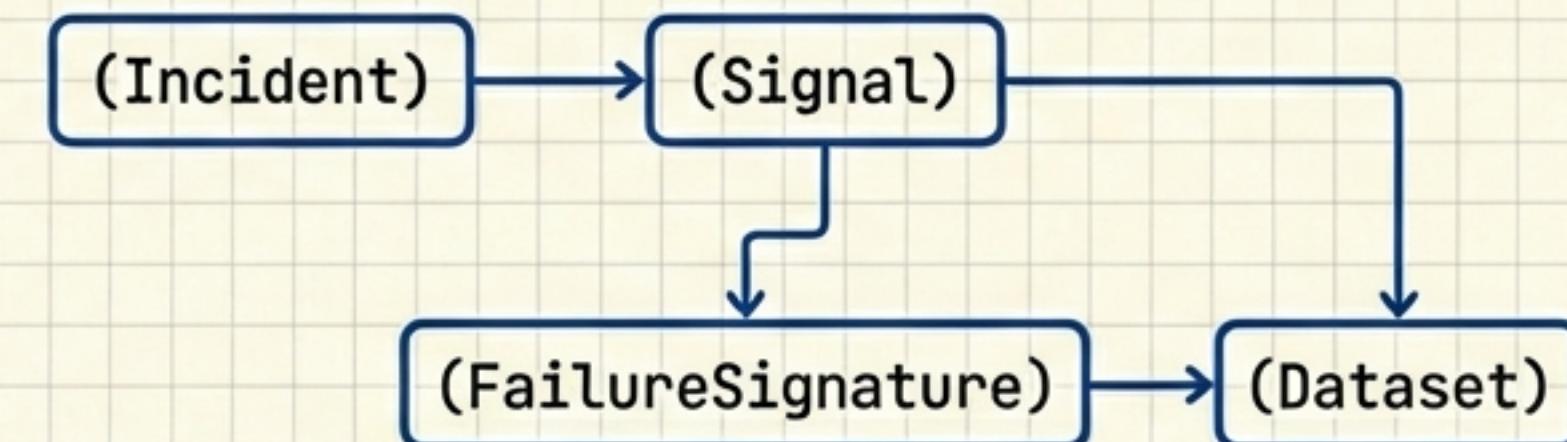
DynamoDB (Operational State)

What is happening right now?



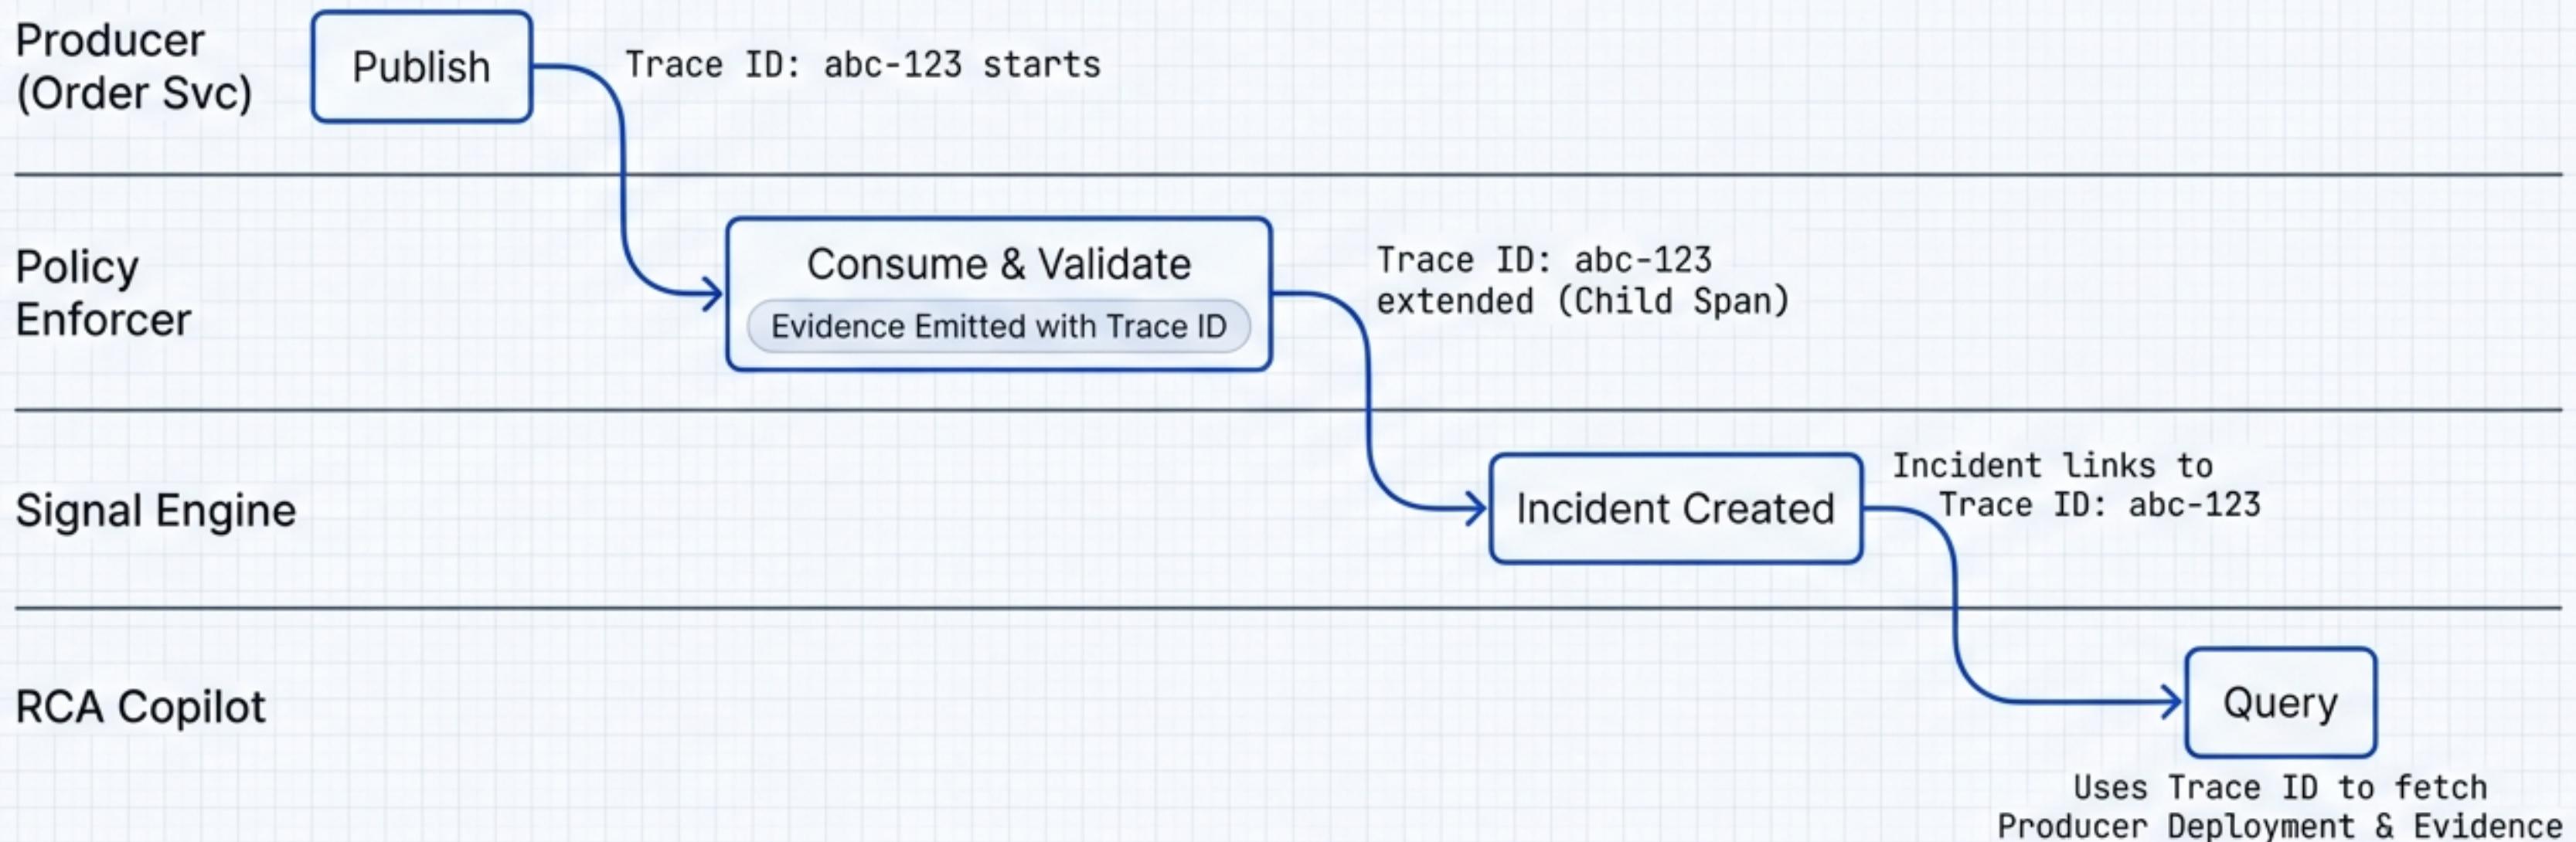
Neptune (Causal Graph)

Why did it happen?



Do NOT write every record to Neptune.
Use **FailureSignature** nodes to compress millions of failures.

Trace Propagation: The Causal Spine



Pattern 1 Implementation: If Trace exists in Raw, extend it. If not, start it.
Always anchor Evidence to Trace.

Failure Modes & Reliability

Scenario	Impact	Mitigation
Enforcer Down	Evidence Gap (Blind spot)	Signal Engine detects 'No Evidence' -> Triggers <code>OBSERVABILITY_PIPELINE_DOWN</code> incident.
Schema Registry Down	Cannot validate schemas	Emit Evidence with reason: <code>SCHEMA_REGISTRY_UNAVAILABLE</code> . Fail open or closed based on policy.
Consumer Lag	Stale Signals	HPA on CPU + Lag metrics. Signal Engines must be lag-aware.

Implementation Blueprint (SRE View)

1 Deployment Stack

- ✓ EKS Managed
- ✓ HPA (Scale on Lag + CPU)
- ✓ Multi-AZ Distribution

2 Service Level Objectives (SLOs)

Completeness: **99.99%** raw records generate evidence.

Latency: Evidence emitted < 2s from ingest.

3 Cost & Storage

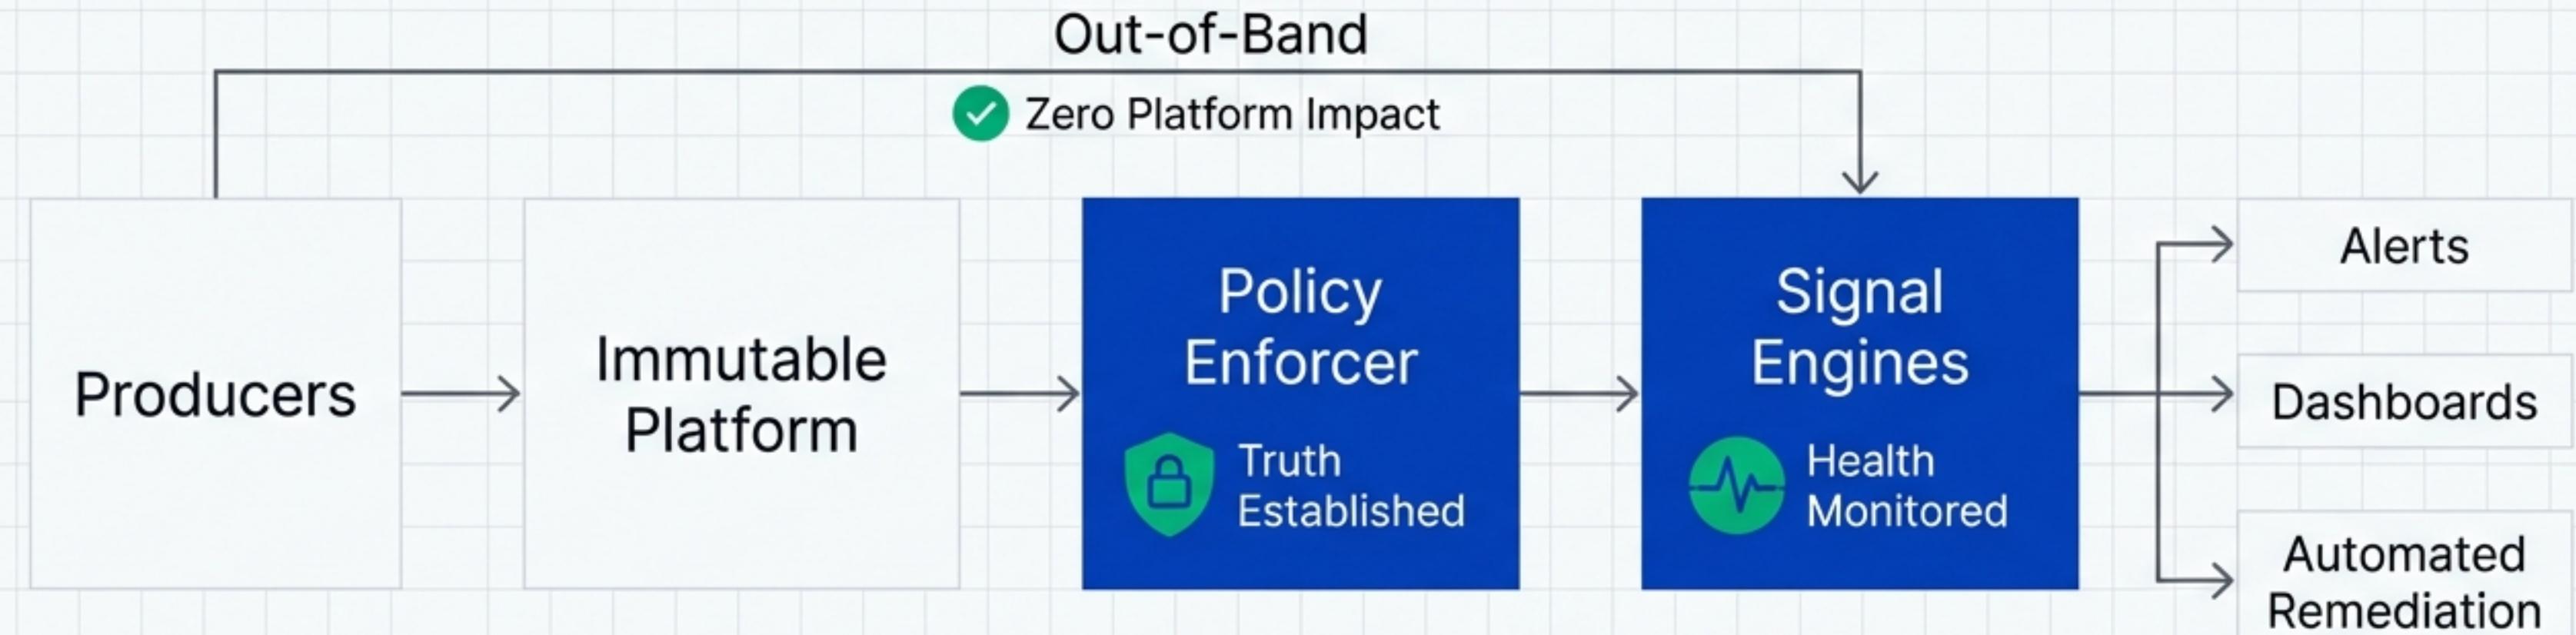
DynamoDB: TTLs enabled (7 days).

Neptune: Async writes.
Sampled/Failures only.

4 Definition of Done

- ✓ Evidence schema versioned.
- ✓ Trace keys present.
- ✓ Signal Engines trusting evidence.

Summary: Immutable Truth, Decoupled Action



We achieved rigorous observability and governance without modifying the immutable platform or producers. The Policy Enforcer provides the Immutable Truth that powers the entire pipeline.