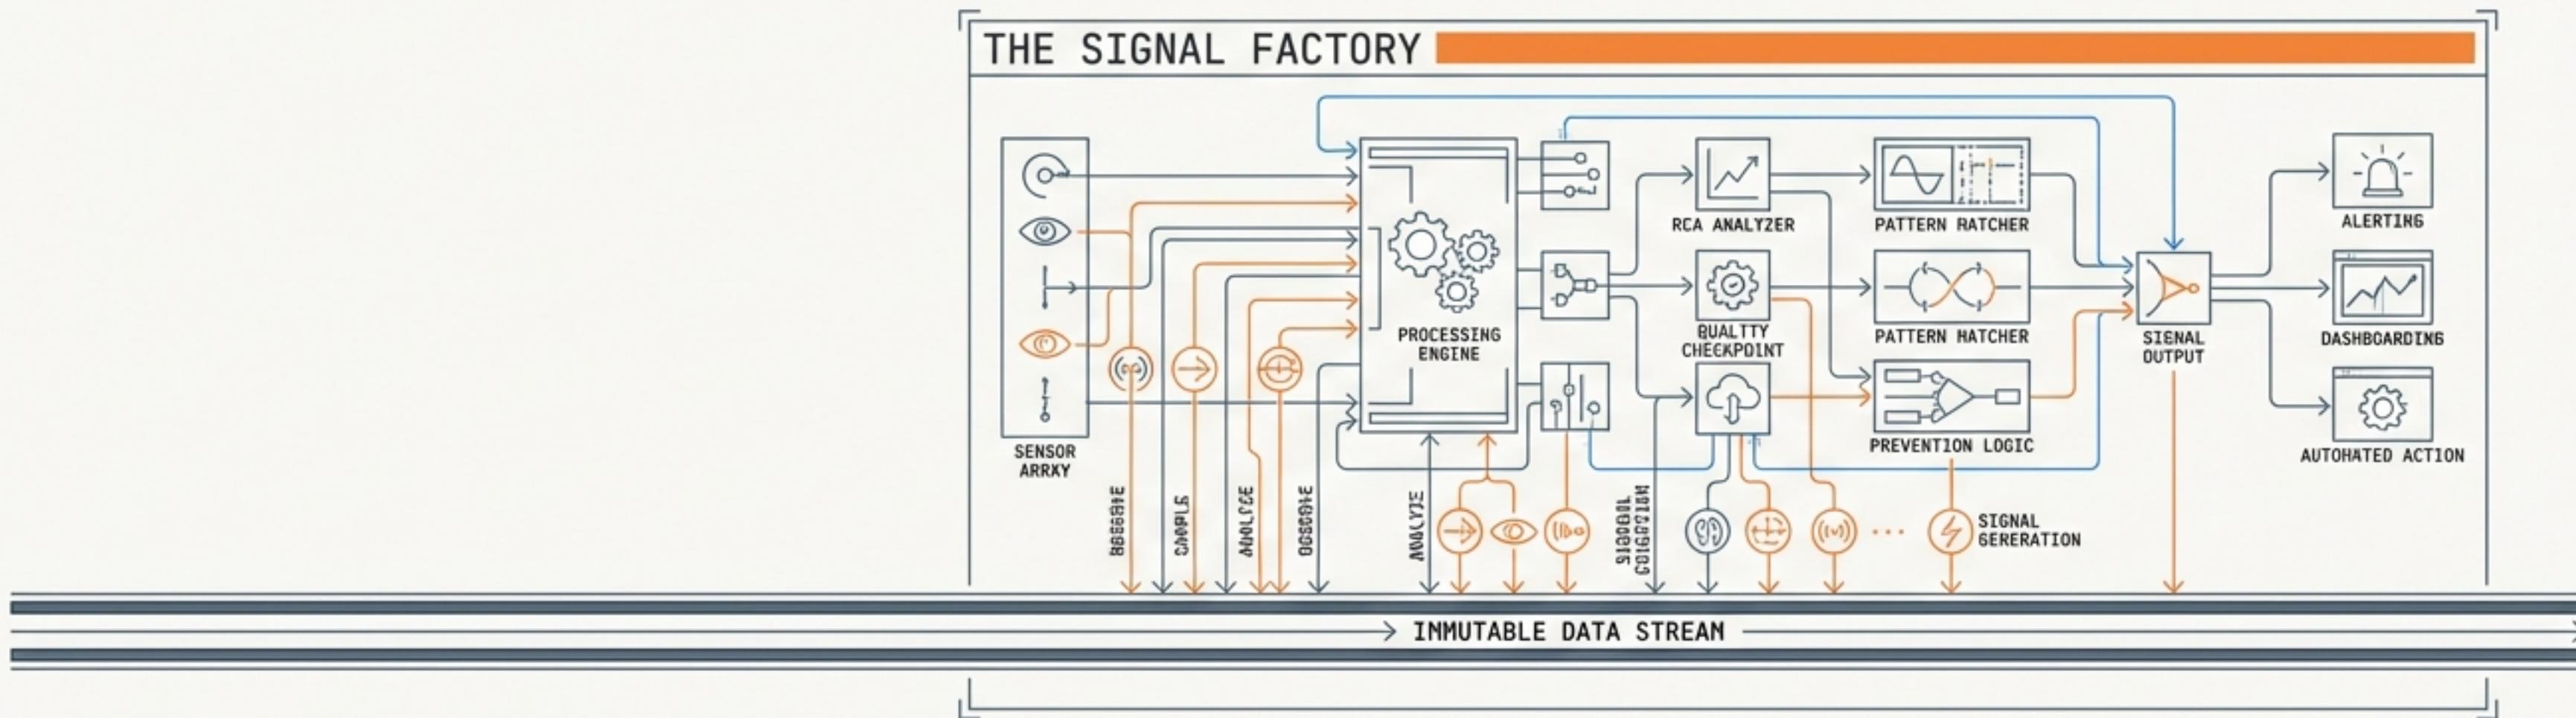


Out-of-Band Observability: The 'Signal Factory' Architecture

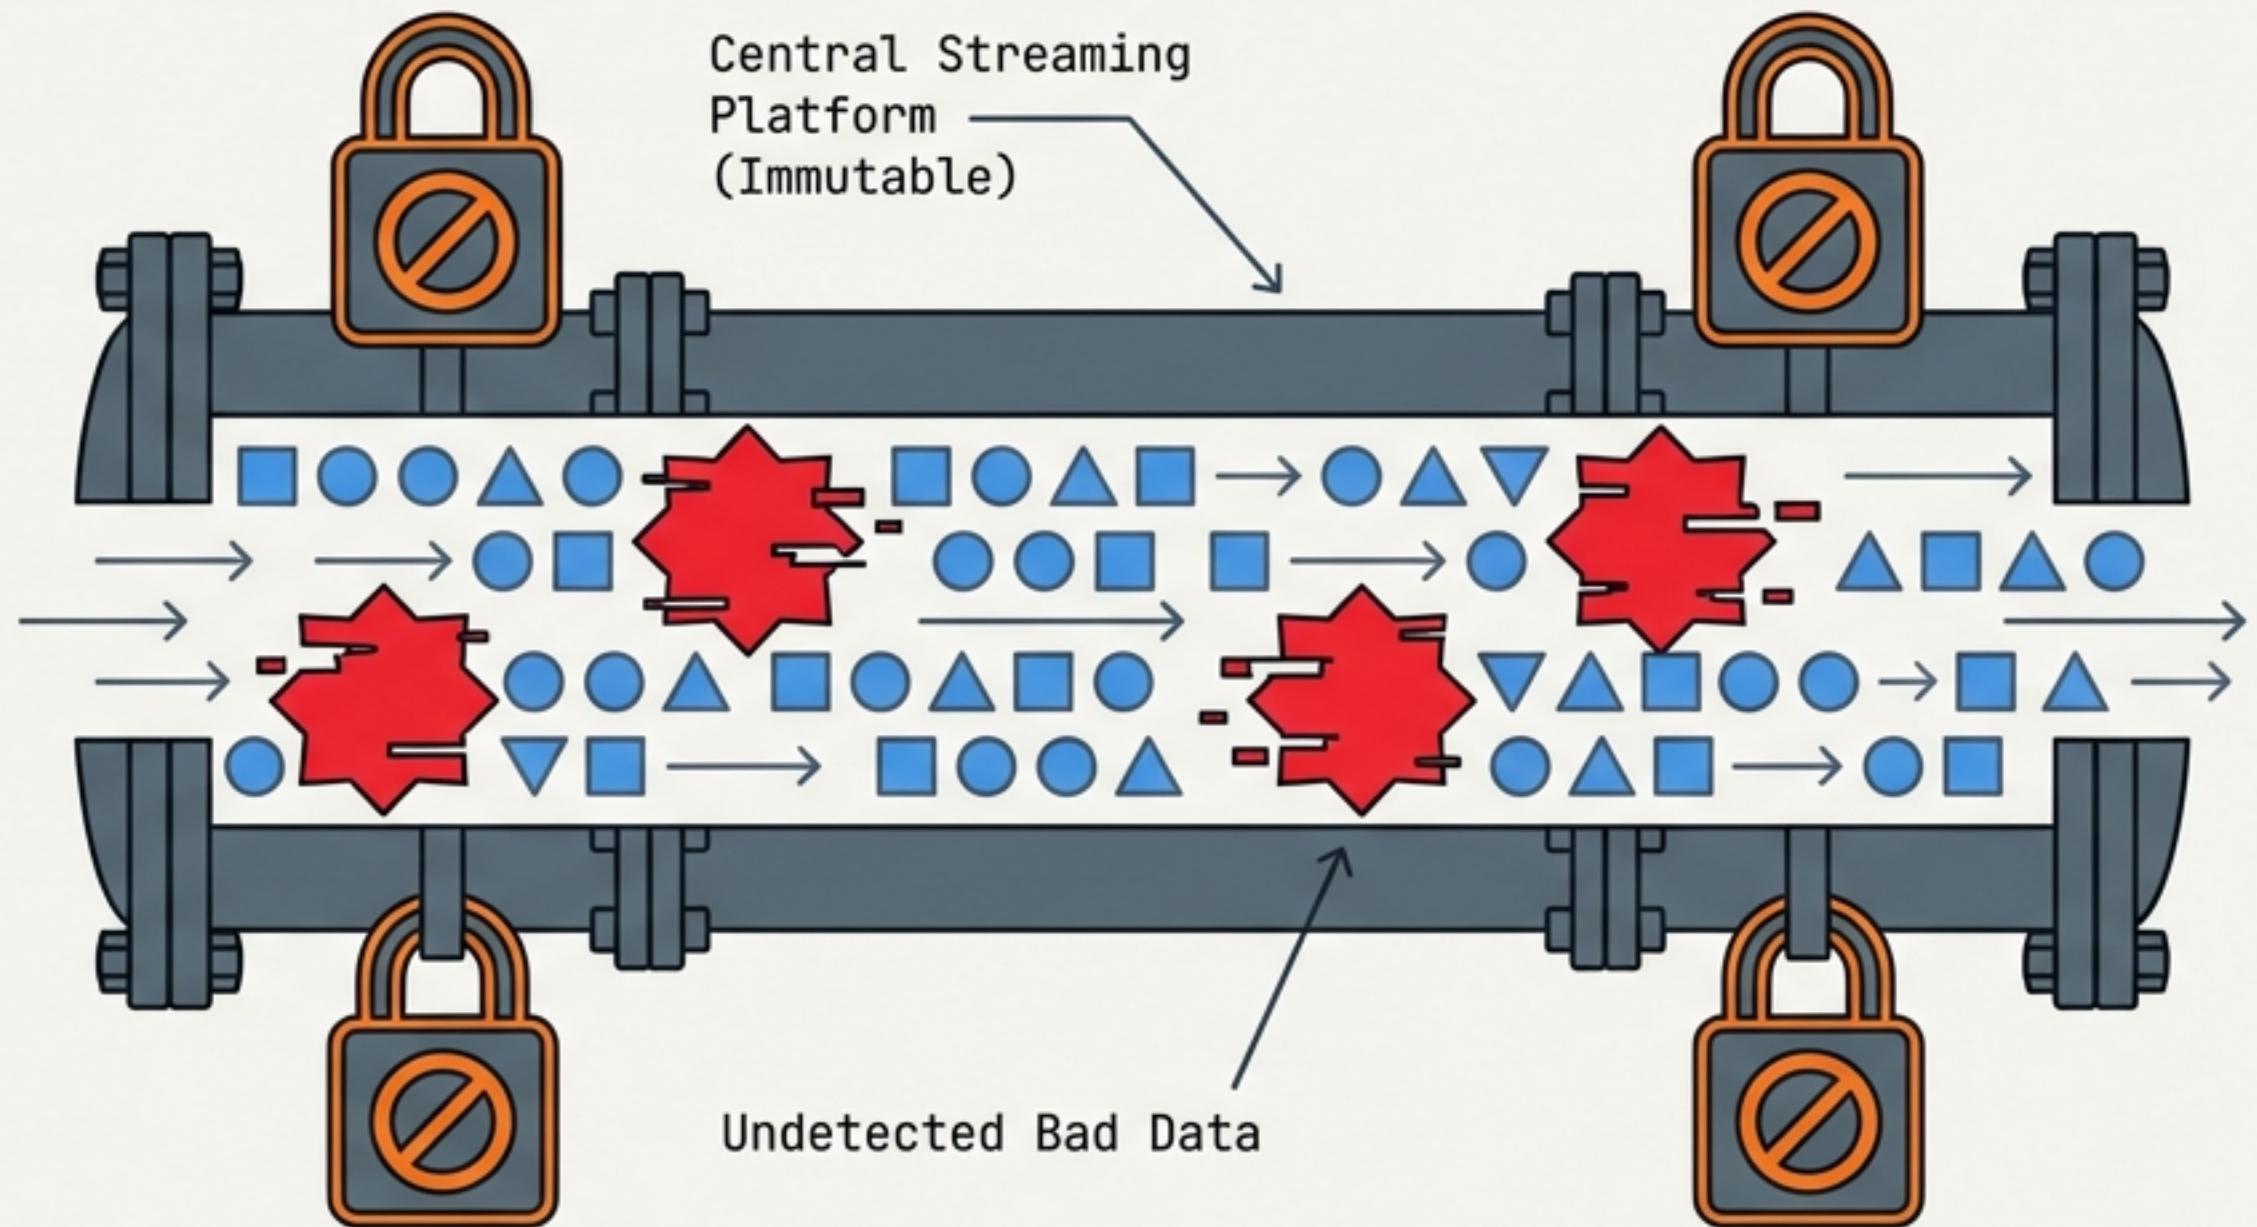
Achieving Prevention, Quality, and Root Cause Analysis (RCA)
without Modifying Producers or the Central Platform.



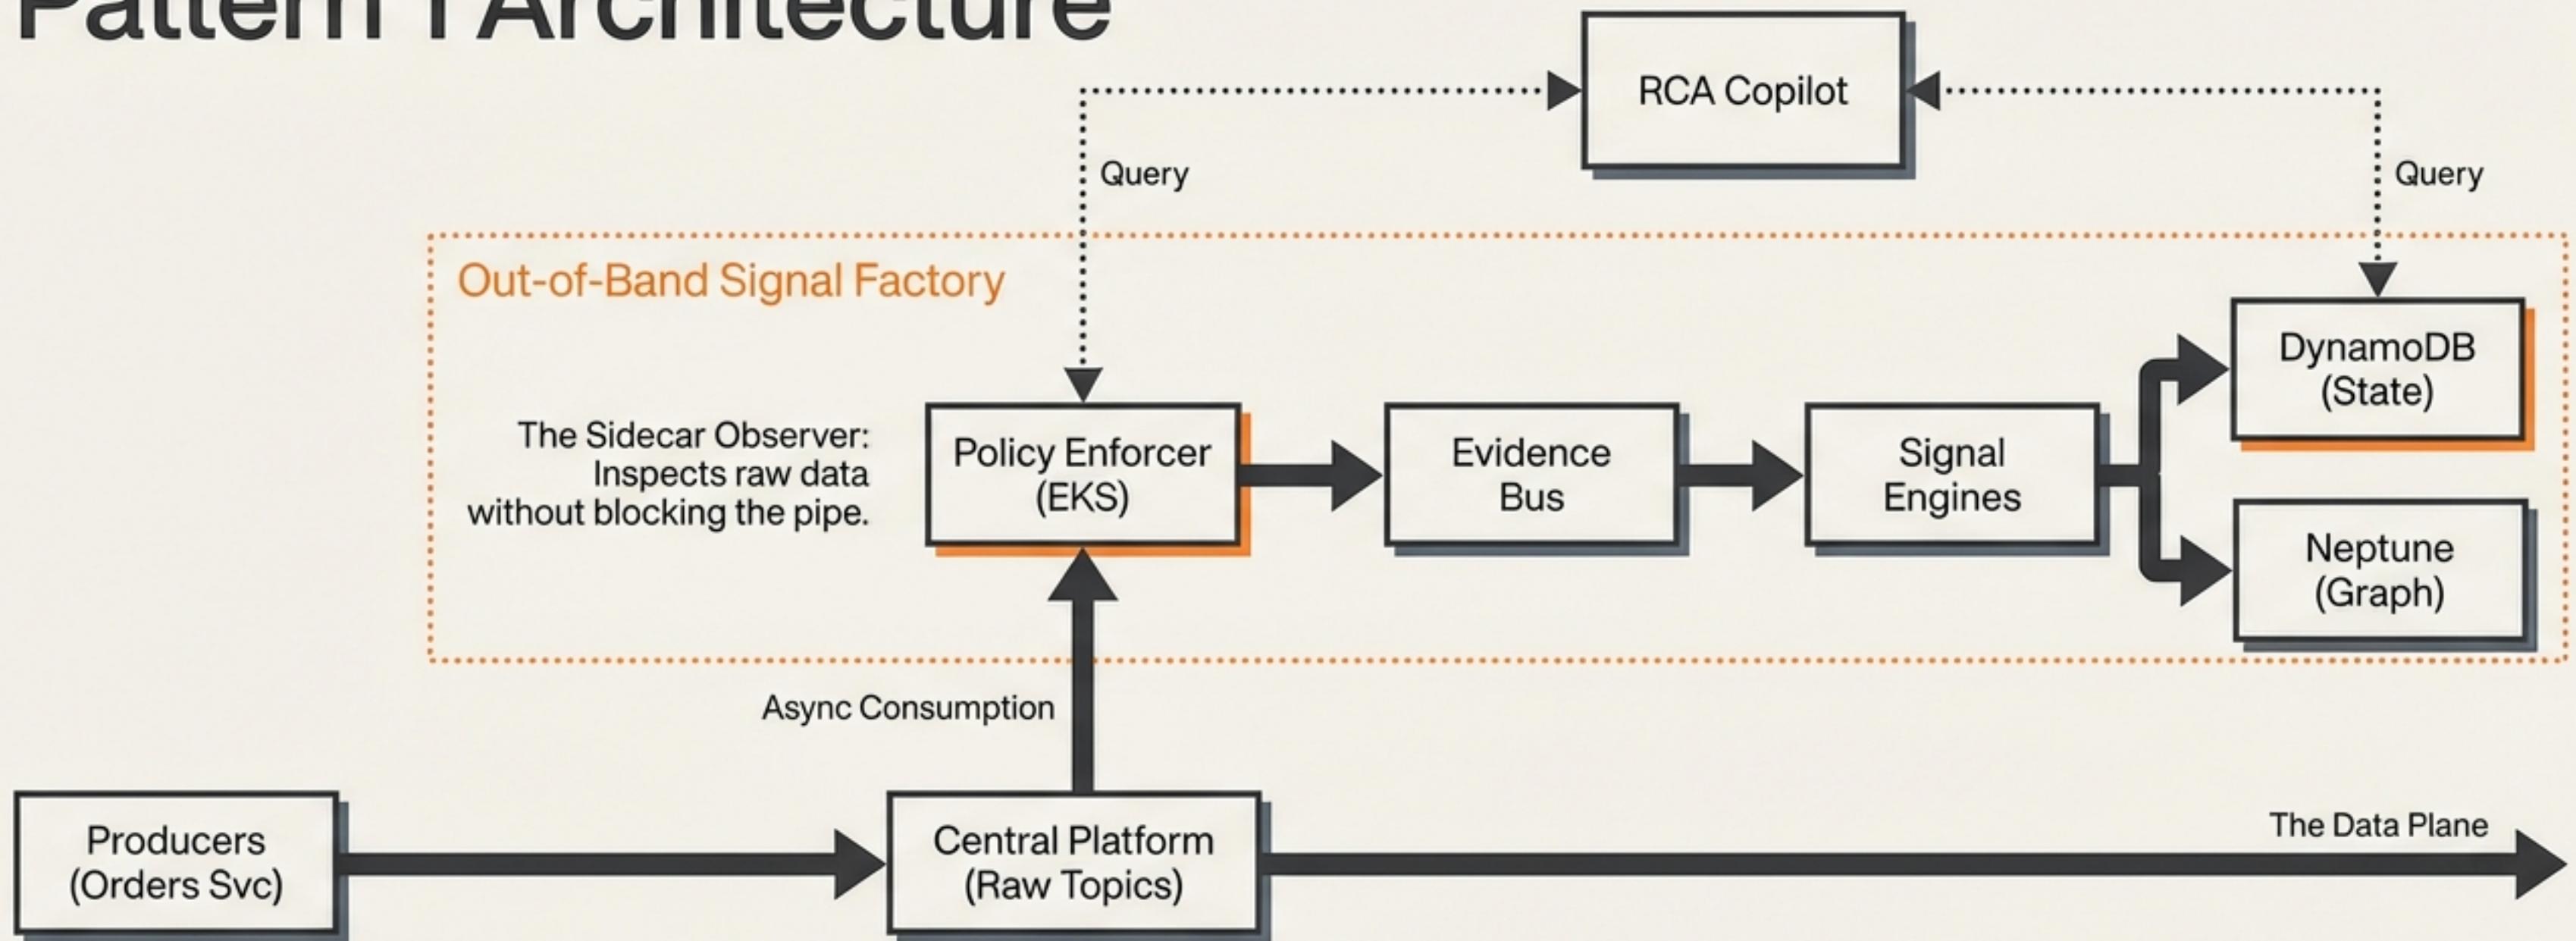
The Challenge: Observability in an Immutable World

We operate in a constrained environment. The **Central Streaming Platform (Kafka/MSK)** acts as a **rigid, immutable pipe**. Producers stream business events directly into this pipe and cannot be modified to add observability SDKs. The Platform itself allows no ingress hooks or routing interception.

The Risk: Bad data (schema violations, PII leaks) flows undetected to consumers until downstream systems crash.



High-Level Solution: Pattern 1 Architecture



The Core Separation: Record Truth vs. System Health



Policy Enforcer

ESTABLISHING RECORD TRUTH

- **Scope:** Atomic. (One record at a time).
- **Responsibility:** Deterministic Validation.
- **Checks:** Schema, Contract, PII.
- **Output:** Immutable Evidence (PASS/FAIL).
- **Key Question:** “Is this specific record valid?”



Signal Engines

ESTABLISHING SYSTEM HEALTH

- **Scope:** Aggregated. (Time windows & datasets).
- **Responsibility:** Anomaly Detection.
- **Checks:** Thresholds, SLOs, Drift.
- **Output:** Signals & Incidents (SEV-1).
- **Key Question:** “Is the system healthy?”

The Journey Begins: A Flawed Event Enters the Stream

Scenario:

The Orders Service deploys a hotfix at 09:58 AM.

The Flaw:

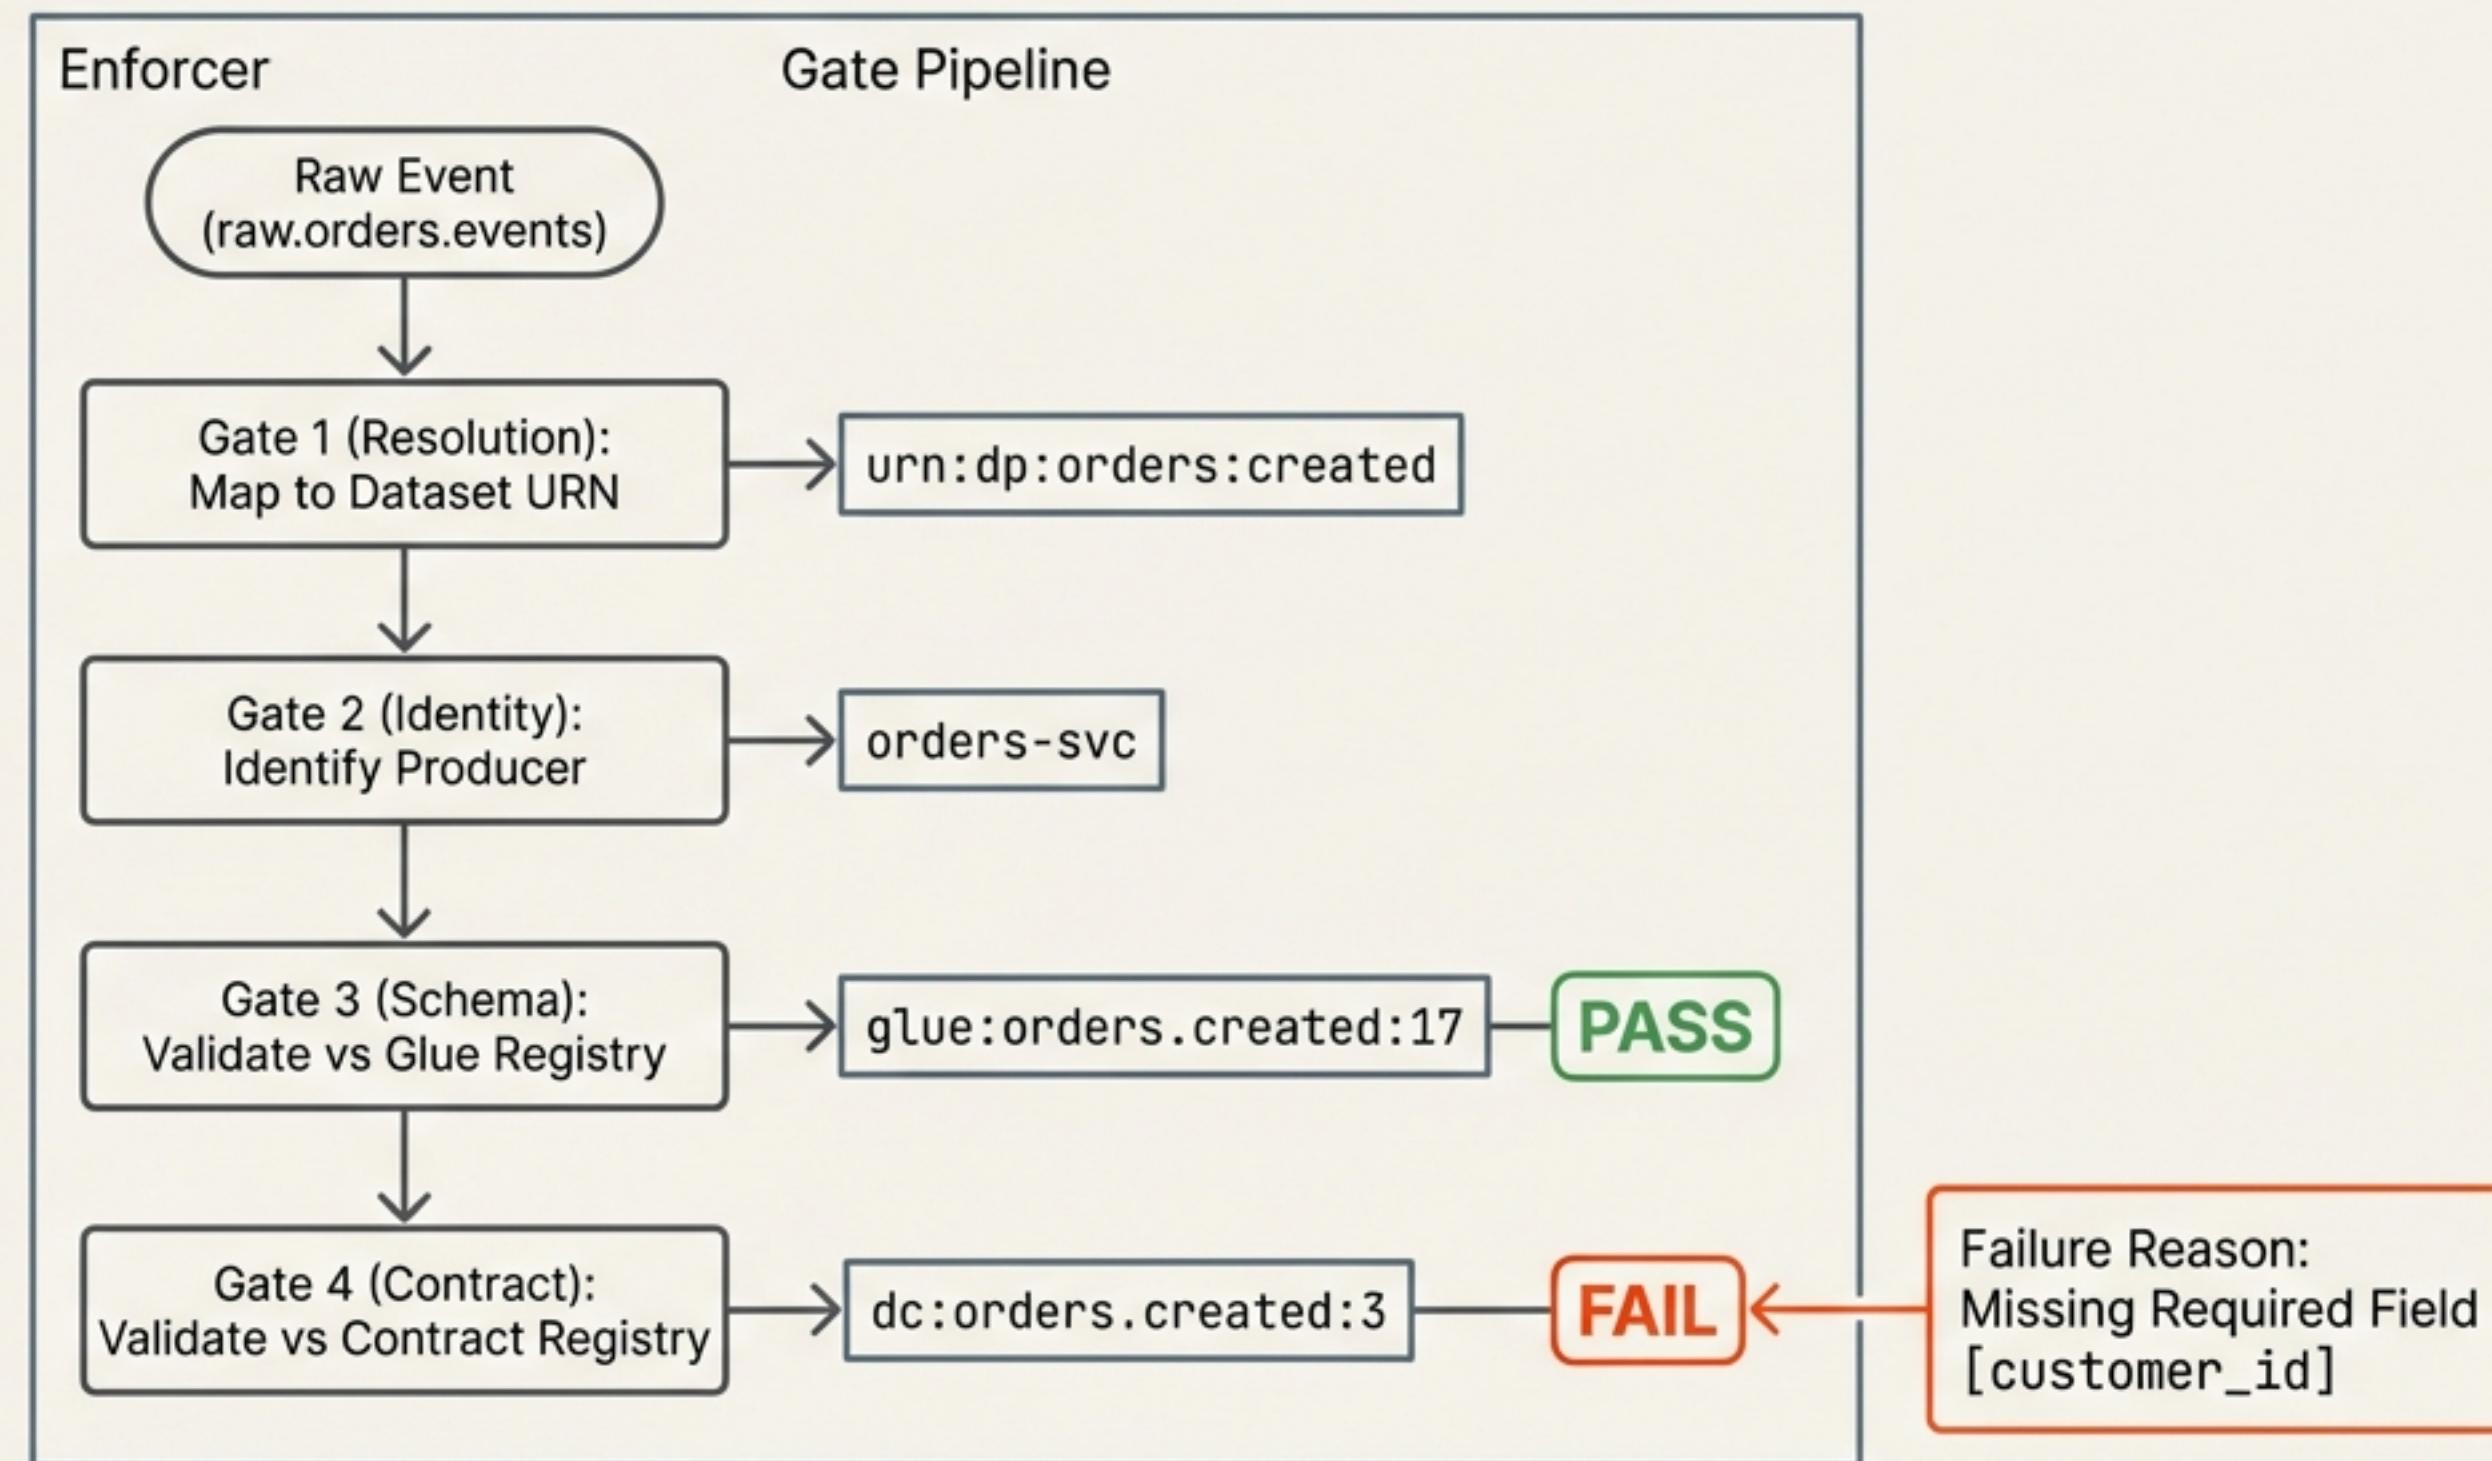
The payload is missing the mandatory `customer_id` field.

Context:

The Central Platform accepts this message blindly because it is immutable.
The bad data is now in the pipe.

```
{  
  "topic": "raw.orders.events",  
  "partition": 4,  
  "offset": 9918273,  
  "payload": {  
    "order_id": "0-99123",  
    "timestamp": "2023-10-27T09:58:01Z",  
    "total_amount": 149.99,  
    "currency": "USD",  
    "items": [  
      {"sku": "SKU-11", "qty": 1}  
    ]  
  }  
}  
← // MISSING: customer_id
```

Step 2: The Policy Enforcer Establishes Truth



The Enforcer does not page on-call yet. It simply establishes the FACT of the failure.

Step 3: Canonical Evidence (The Interface)

This immutable object IS the observability product.

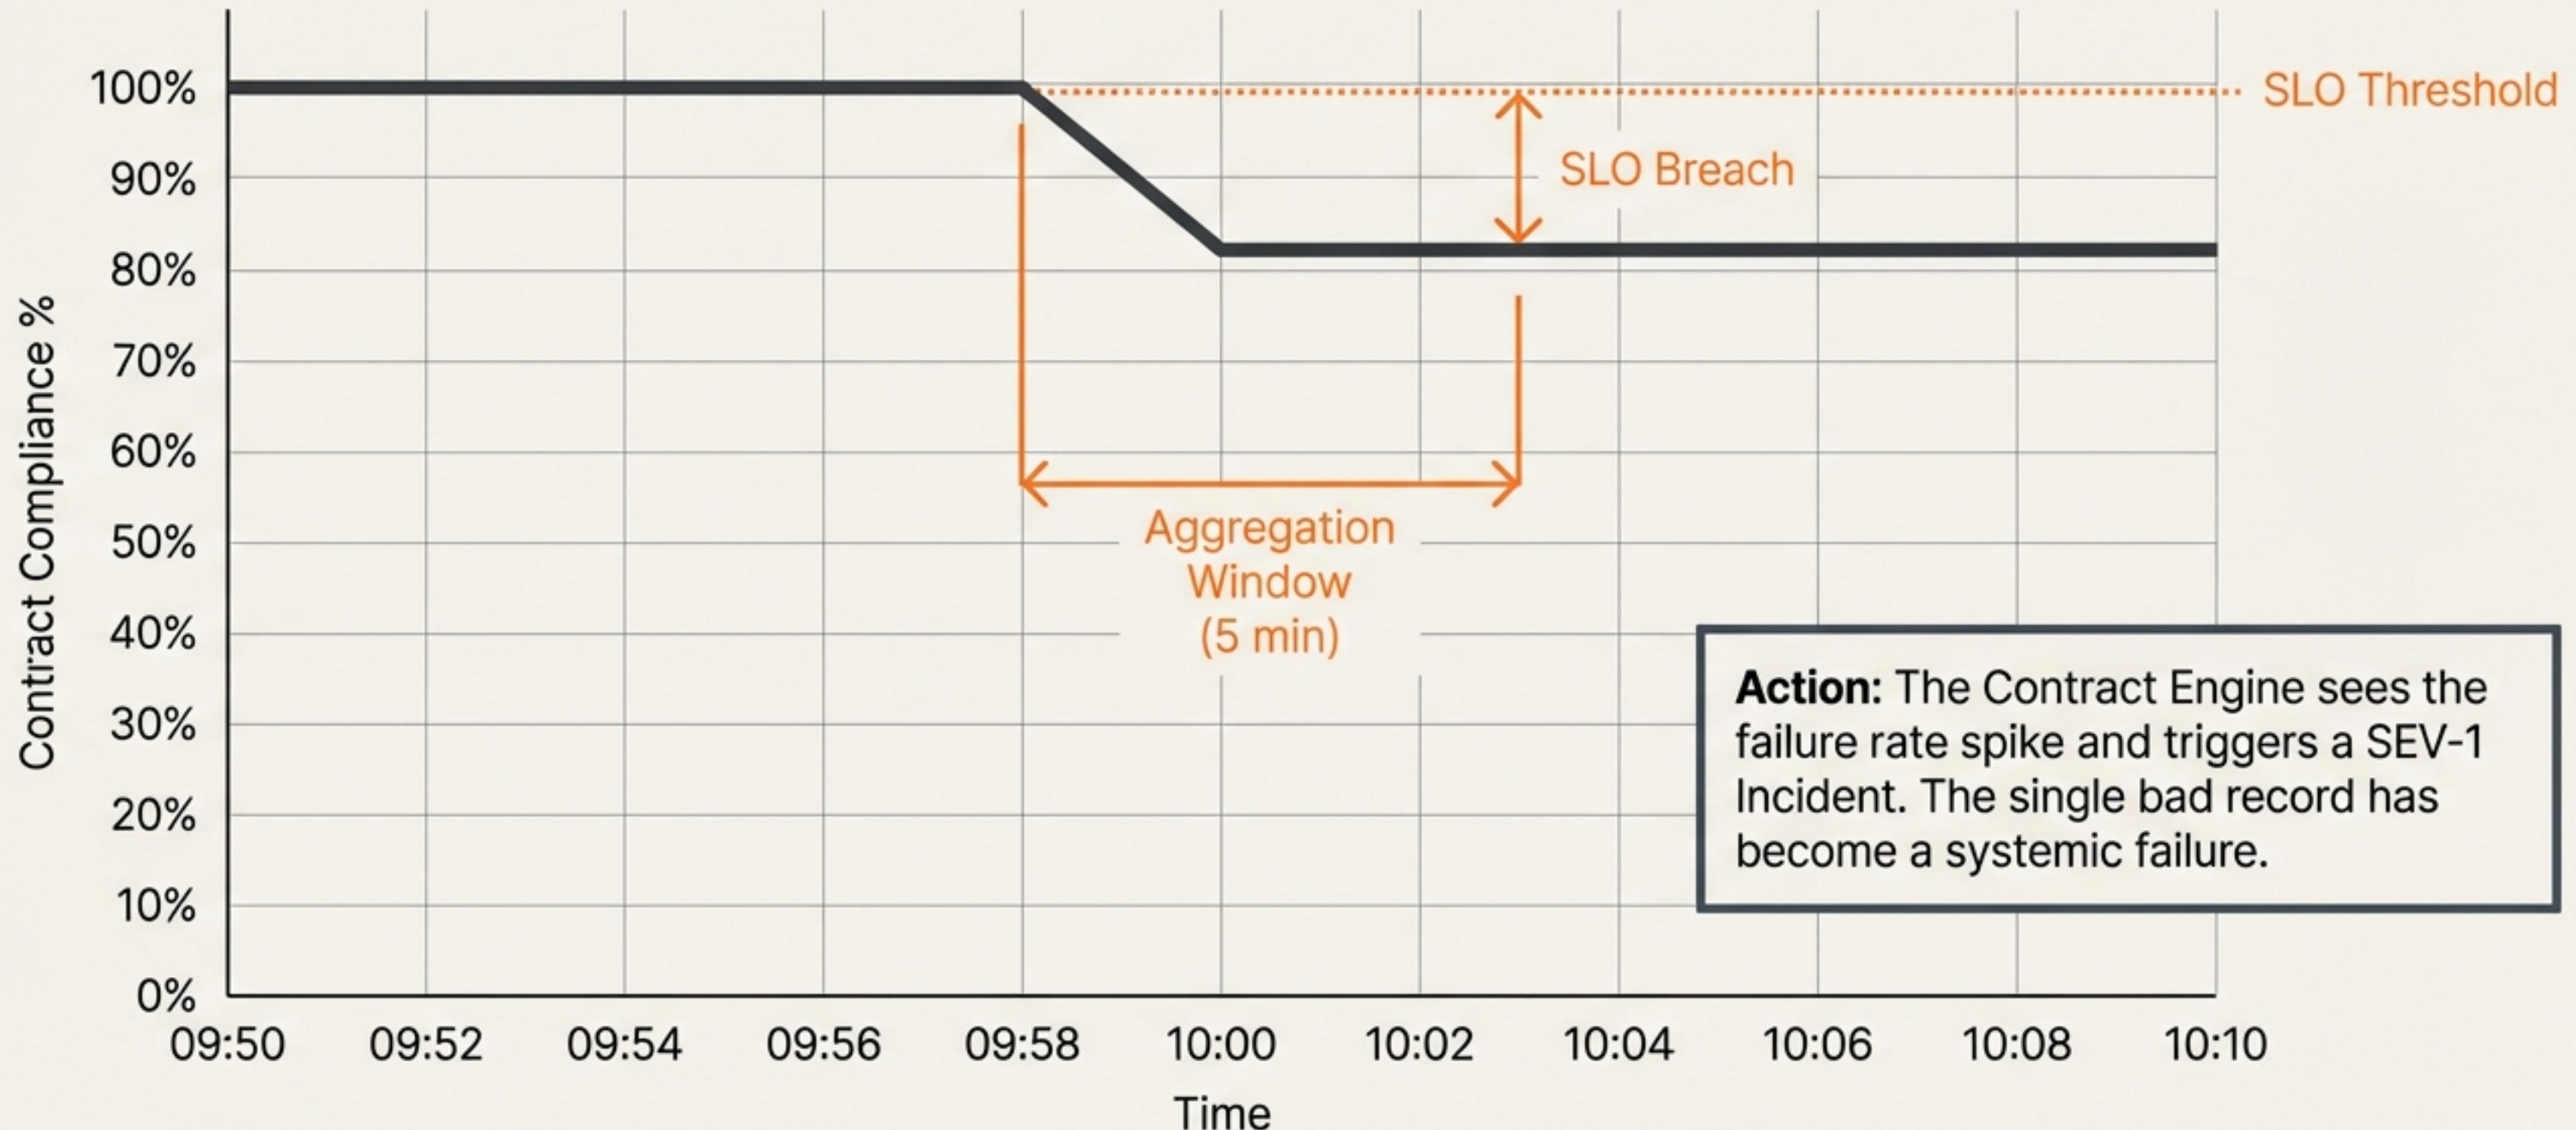
```
{  
  "evidence_id": "evd-01J...",  
  "timestamp": "2023-10-27T09:58:02Z",  
  "dataset_urn": "urn:dp:orders:created",  
  "producer": {  
    "id": "orders-svc",  
    "confidence": "HIGH"  
  },  
  "validation": {  
    "result": "FAIL",  
    "failed_gates": ["CONTRACT"],  
    "reason_codes": ["MISSING_FIELD:customer_id"]  
  },  
  "source": {  
    "topic": "raw.orders.events",  
    "offset": 1882341  
  },  
  "otel": {  
    "trace_id": "ab91f..."  
  }  
}
```

The Logical Context →

The Deterministic Fact

The Join Key (Replayability)

Step 4: Signal Processing Engines Establish Health



Step 5: Incident Creation & Trace Anchoring

Incident Ticket

Title: SEV-1 Contract Violation

Service: Orders Service

ID: INC-7721

Primary Trace ID: ab91f...

Trace Chain

Producer: orders-svc

Topic: raw.orders.events

✖ Enforcer Span: VALIDATION_FAIL

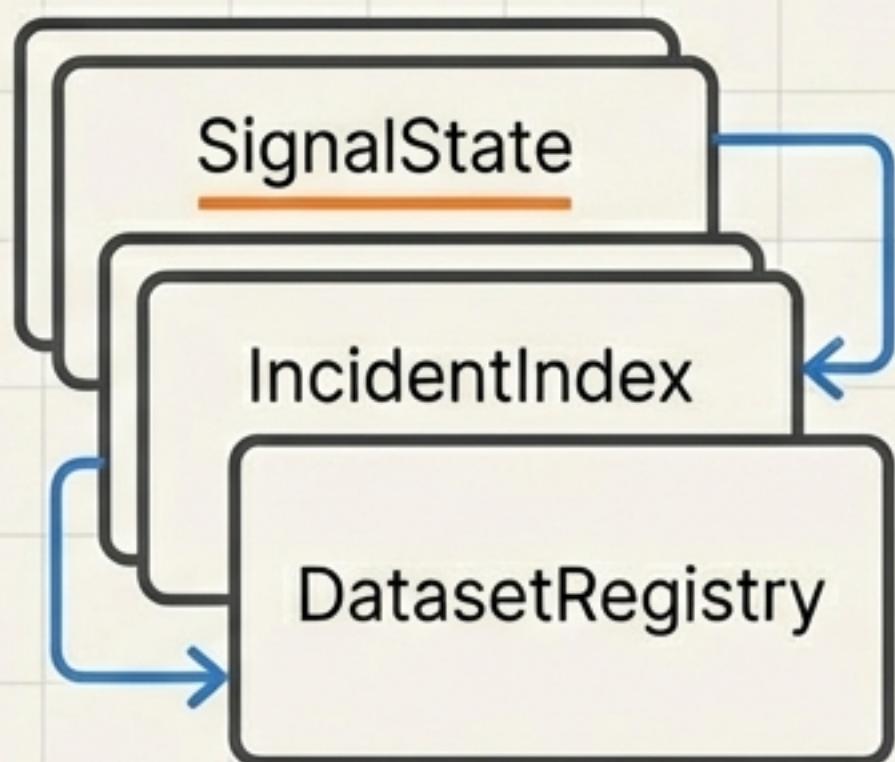
Evidence ID: evd-01J...

Value: The Incident is grounded in the Trace ID. One click takes the SRE from the alert to the exact sequence of events.

Data Store Architecture: State vs. Causality

DynamoDB

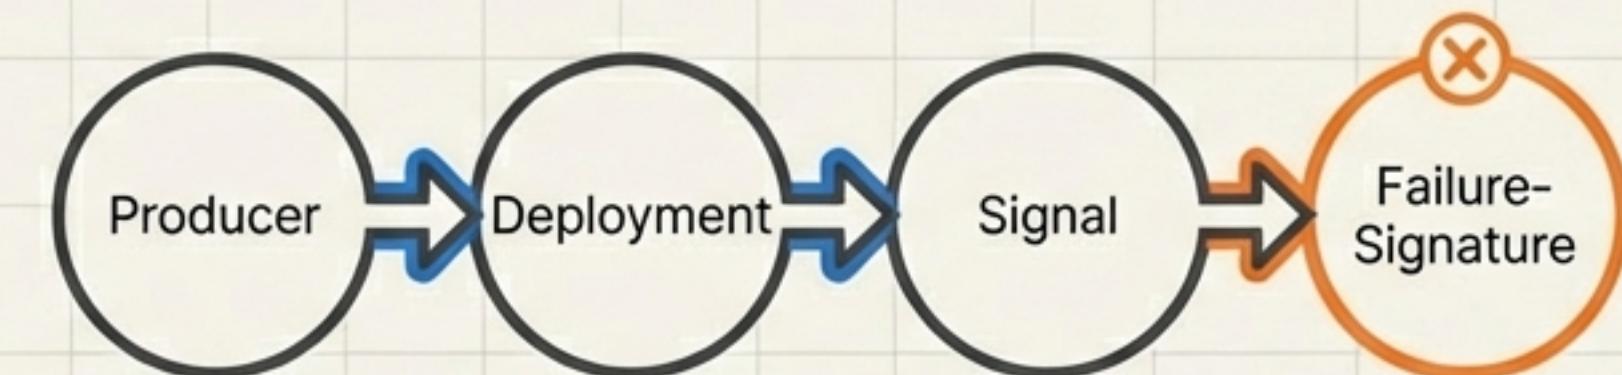
Operational Truth (State)



Use case: Fast Dashboards.
“What is the current health?”

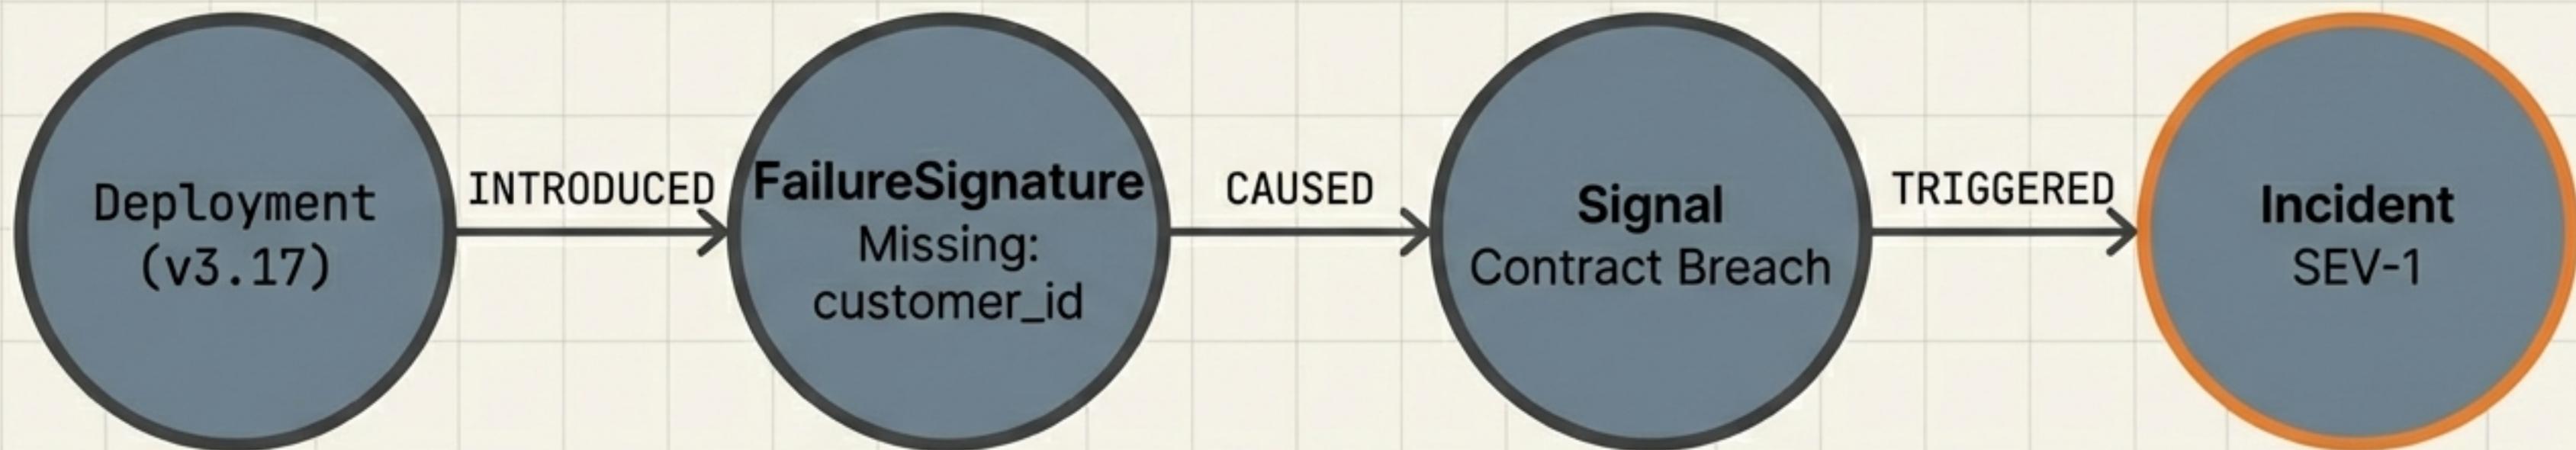
Amazon Neptune

Causal Graph (Relationships)



Use case: Root Cause Analysis.
“Why did this happen?”

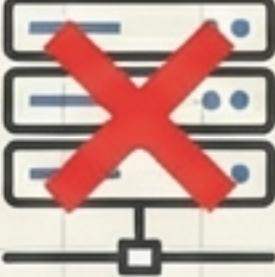
The Graph Model: Automatic RCA



Optimization: We do not store every failed record as a node. We bucket them into unique Failure Signatures to prevent graph explosion.

Resilience & Failure Modes

Observability of Observability

| Enforcer Down | Impact | Action |
|---|---------------------------------------|---|
|  | Evidence Gap. | <p>Signal Engines detect "Zero Evidence" vs Raw Lag. Trigger `OBSERVABILITY_PIPELINE_DOWN`.</p> |
|  | Impact Cannot validate schema. | <p>Action Enforcer falls back to WARN mode. Emits evidence with `failed_gates=["REGISTRY_UNAVAILABLE"]`. Never crash the consumer.</p> |
|  | Impact Delayed signals. | <p>Action HPA scales consumers. Signal Engine annotations mark data as "Stale".</p> |

Summary of Responsibilities

| Responsibility | Owner | Artifact |
|-----------------------|------------------------------|-----------------------|
| Record Validity | Policy Enforcer | Evidence (PASS/FAIL) |
| System Health | Signal Engines | Signals & Incidents |
| Alert Thresholds | Signal Engines | Dataset Policy Config |
| Schema Bindings | Gateway Control Plane | Registry Map |
| RCA Narrative | RCA Copilot | Causal Graph Path |

The Result: Prevention-Grade Observability

Zero

Producer Changes Required

The Central Platform remains immutable. Enforcer runs entirely out-of-band.

< 2s

Latency

From raw ingest to Evidence emission and Signal update.

100%

Attribution

Failures are linked to producers via best-effort identity and confidence scores.

By moving enforcement to the sidecar, we gain the safety of a gateway with the flexibility of a decoupled observer.